



第九章 空间网络模型与查询

陶煜波
计算机科学与技术学院

第八章 空间查询处理与优化回顾

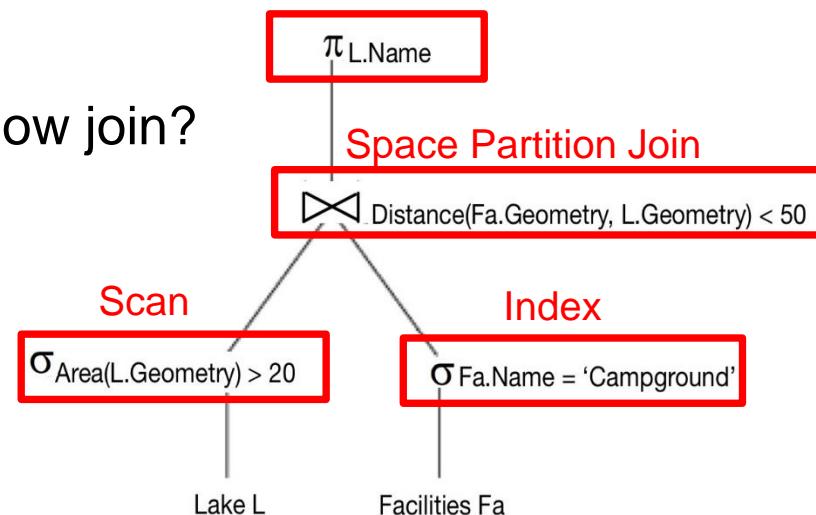
- Query processing and optimization (QPO)
 - Translates SQL Queries to execution plan
- Three key concepts in QPO
 - Building blocks
 - Decompose SQL query into building blocks, e.g., select, join, sort
 - Query processing strategies for building blocks
 - A few processing strategies for each building block
 - E.g. a point query – scan vs. index
 - Query optimization
 - For each building block of a given query, DBMS QPO tries to choose “most efficient” strategy given database parameters
 - A fixed priority scheme
 - Rule based approach
 - Cost-based model

第八章 空间查询处理与优化回顾

- Building blocks and strategies
 - Point Query
 - Linear search, binary search with Z-Curve, Indexed search with R-Tree
 - Range Query
 - Linear search, binary search (+scan) with Z-Curve, Indexed search with R-Tree
 - Nearest Neighbor
 - Two phase approach, single phase approach
 - Spatial Join
 - Nested loop, nested loop with one spatial index, space partitioning, tree matching

第八章 空间查询处理与优化回顾

- QPO process steps include
 - Creation of a query tree for the SQL query
 - Choice of strategies to process each node in query tree
 - Ordering the nodes for execution
- Key ideas for SDBMS include
 - Filter-Refine paradigm to reduce complexity
 - New building blocks and strategies for spatial queries
 - CPU cost is higher
 - Push down spatial selection below join?
 - Reorder join operations



几何对象模型

- 空间数据模型分类

- 矢量模型

- 几何对象模型 (地物几何形状)

- 几何拓扑模型 (地物几何形状 + 拓扑)

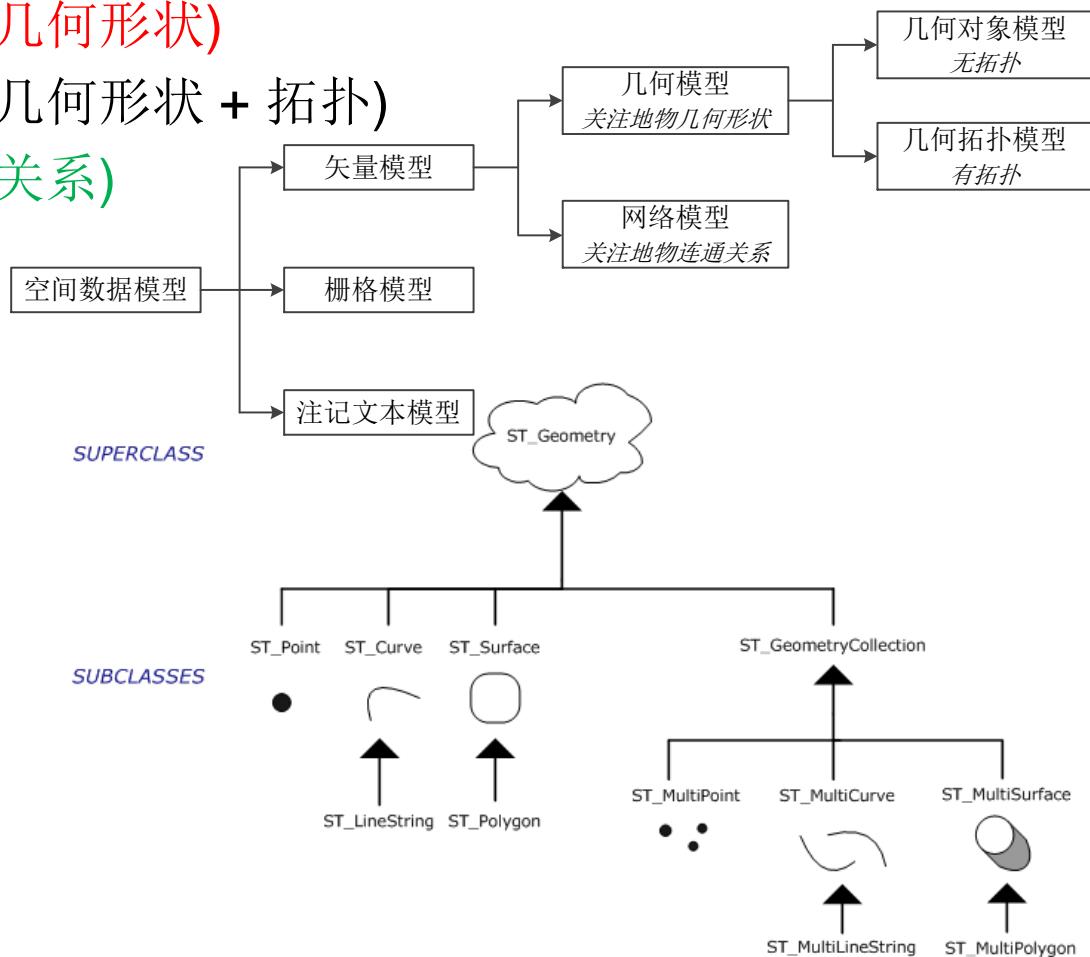
- 网络模型 (地物连通关系)

- 栅格模型

- 注记文本模型

- 几何对象模型

- 数据 + 方法



空间数据库

- 空间数据库 = 对象关系/关系数据库 + 空间扩展
 - Oracle + Oracle Spatial
 - SQL Server + SQL Server Spatial
 - PostgreSQL + PostGIS
 - MySQL + MySQL Spatial
 - SQLite + SQLite Spatialite
- PostGIS
 - 提供了空间数据类型、空间函数和空间索引
 - geom geometry (Point/Line/Polygon/Multixxx, 空间参考系)
 - ST_GeomFromText ('WKT表示', 空间参考系)
 - ST_XXX: ST_Distance, ST_Intersects, ST_DWithin, ...
 - GiST

应用举例：打车软件

- 通常，数据库保留所有时间上的信息，即
 - Taxi(ID, driverID,, status, pos(Point, 4326), time)
 - User(ID, name,, pos(Point, 4326), time)
 - Road(ID, name,, line(LineString, 4326))
- 出租车(ID = B)附近1公里内的乘客叫车？
 - Select T.ID, T.position
From Taxi T, User U
Where T.ID = B and
ST_Distance(T.pos, U.pos, true) < 1000
and T.time
and U.time
 - 这样的实现方式合理吗？

第九章 空间网络模型与查询

- 9.1 Motivation and use cases
- 9.2 Conceptual model
- 9.3 Logical model
 - 9.3.1 Transitive Closure
 - 9.3.2 CONNECT statement
 - 9.3.3 RECURSIVE statement
 - 9.3.4 RECURSIVE in PostgreSQL
- 9.4 Physical model
 - 9.4.1 Storage and data structures
 - 9.4.2 Algorithms for connectivity query and shortest path
- 9.5 pgRouting

参考教材：
Spatial Databases: A Tour Chapter 6
空间数据库管理系统概论，3.3-3.4

Navigation Systems

- Historical
 - Navigation is a core human activity for ages!
 - Trade-routes, Routes for Armed-Forces
- Recent Consumer Platforms
 - Devices: Phone Apps, In-vehicle, “GPS”, ...
 - WWW: Google Maps, MapQuest, ...
- Services
 - Display map around current location
 - Compute the shortest route to a destination
 - Help drivers follow selected route



Location Based Services

- Location: Where am I?
 - Geo-code: Place Name (or Street Address) → <latitude, longitude>
 - Reverse Geo-code: <latitude, longitude> → Place Name
- Directory: What is around me?
 - Where is the nearest Clinic? Restaurant? Taxi?
 - List all Banks within 1 mile
- Routes: How do I get there?
 - What is the shortest path to get there?
 - ...

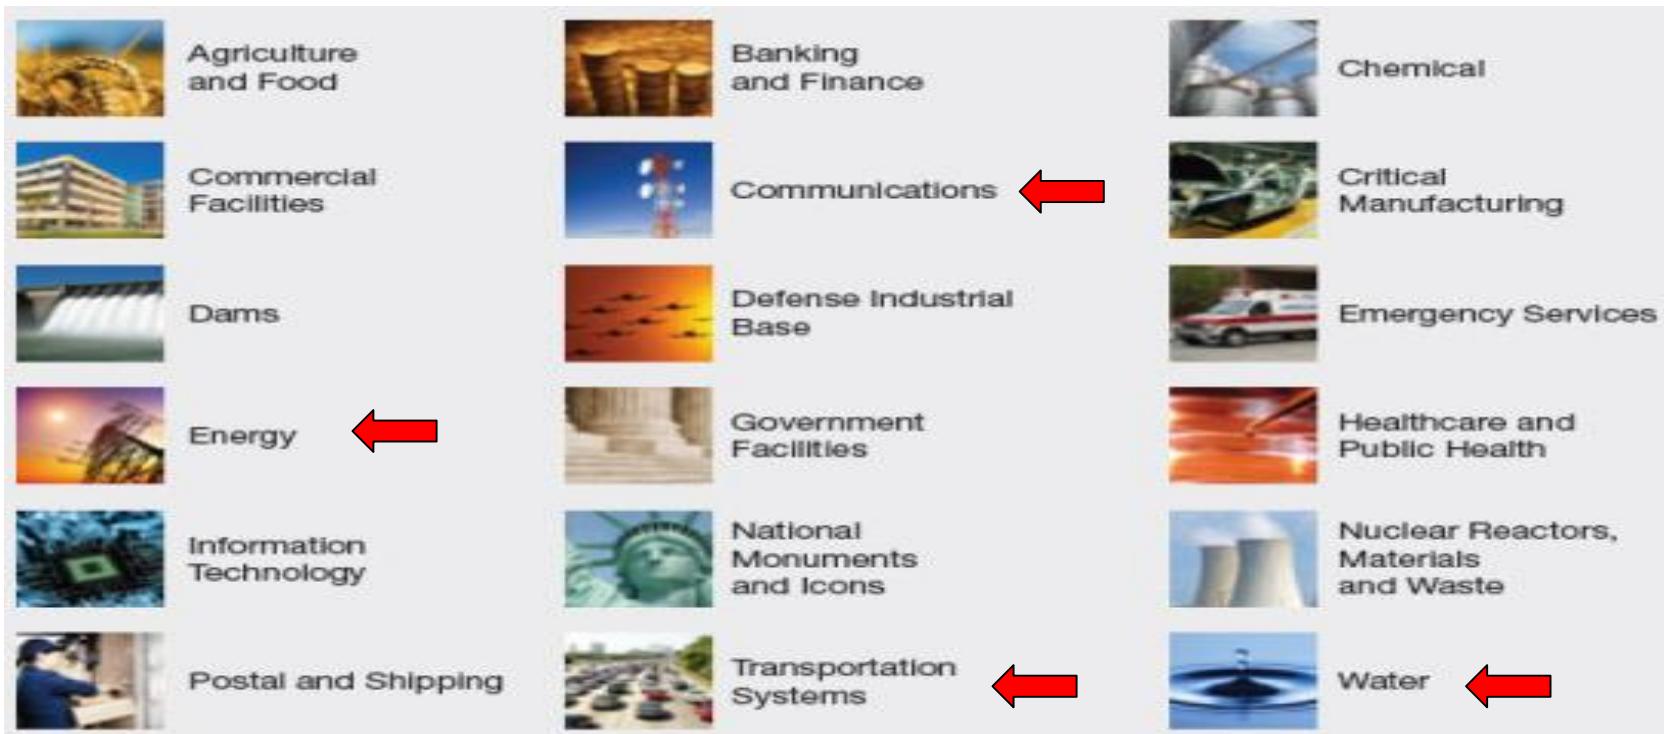
基于位置的服务

- 休闲娱乐型
 - 签到模式
 - 大富翁游戏模式
- 生活服务型
 - 周边生活服务的搜索
 - 与旅游的结合
 - 会员卡与票务模式
- 社交型
 - 地点交友，即时通讯
 - 以地理位置为基础的小型社区
- 商业型 <http://blog.csdn.net/superjunjin/article/details/7818378>
 - LBS+团购；优惠信息推送服务；店内模式

Spatial Networks & Modern Society

- Transportation, Energy, Water, Communications,

...



Source: http://www.dhs.gov/files/programs/gc_1189168948944.shtml

Quiz 1

- Which of the following is not an application of spatial networks?
 - a) Navigation system
 - b) Geocoding
 - c) Reverse geocoding
 - d) Convex hull of a country

Limitations of Spatial Querying

- OGIS Simple Feature Types
 - Supports Geometry (e.g., Points, LineStrings, Polygons, ...)
 - However, lack **Graphs** data type, **shortest_path** operator
- Traditional SQL
 - Supports select, project, join, statistics
 - Lacked transitive closure, e.g., **network analysis** (next slide)
 - SQL3 added **recursion & transitive closure**

思考：transitive closure与函数依赖的闭包之间的关系

Spatial Network Analysis

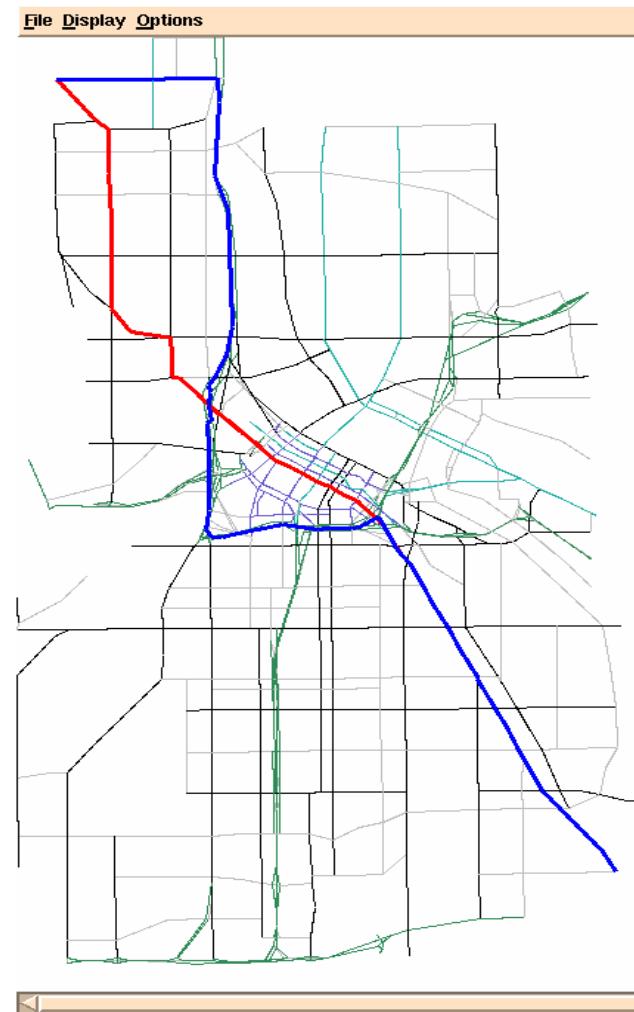
- Route (A start-point, Destination(s))
 - What is the shortest path to get there?
 - What is the shortest path to cover a set of destinations?
- Allocation (A set of service centers, A set of customers)
 - Assign customers to nearest service centers
 - Map service area for each service center
- Site Selection (A set of customers, Number of new service centers)
 - What are best locations for new service centers?

Quiz 2

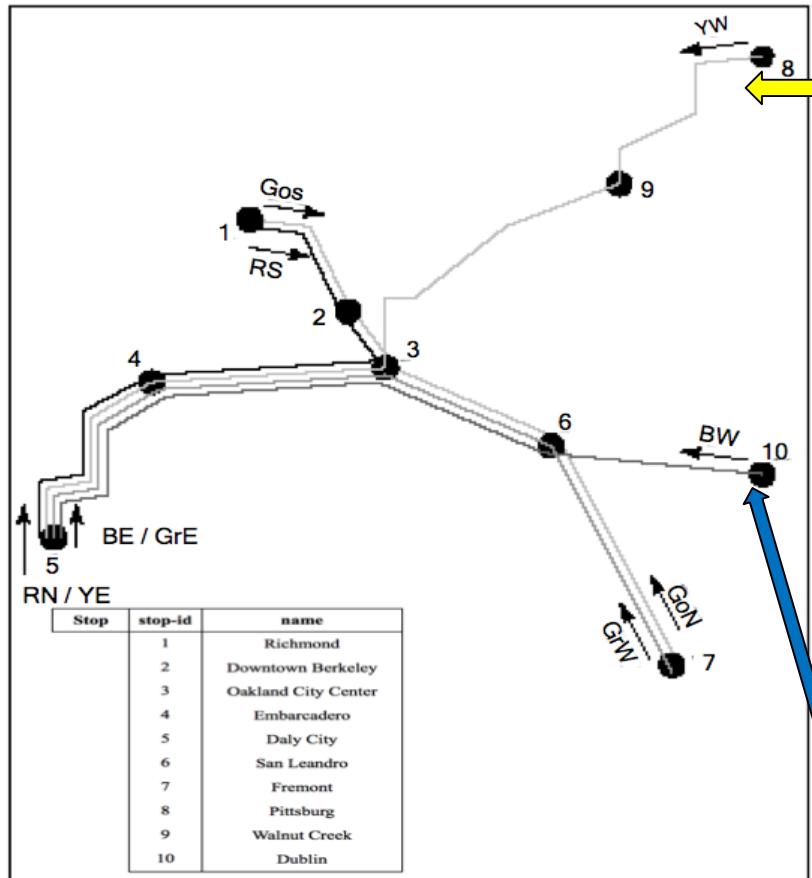
- Which of the following is not included in OGIS simple feature types?
 - a) Graph
 - b) Point
 - c) Line String
 - d) Polygon

Spatial Network Query Example

- Find **shortest path** from a start-point to a destination
- Find **nearest** hospital by driving distance
- Find **shortest route** to deliver packages to a set of homes
- Allocate customers to **nearest** service center



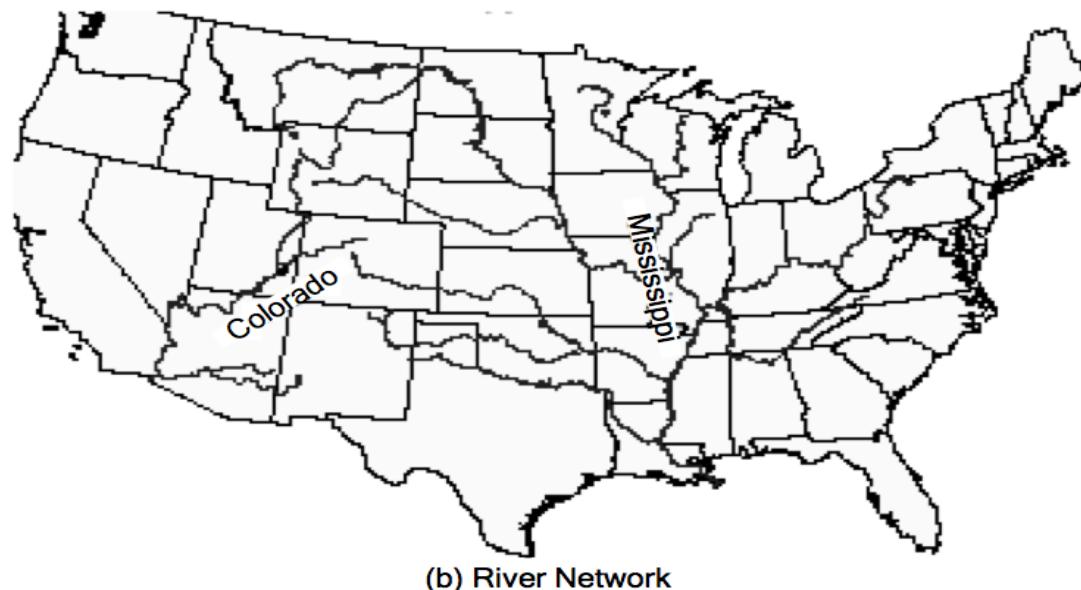
Railway Network & Queries



- Find the number of stops on the Yellow West (YW) route
- List all stops which can be reached from Downtown Berkeley (2)
- List the routes numbers that connect Downtown Berkeley (2) & Daly City (5)
- Find the last stop on the Blue West (BW) route

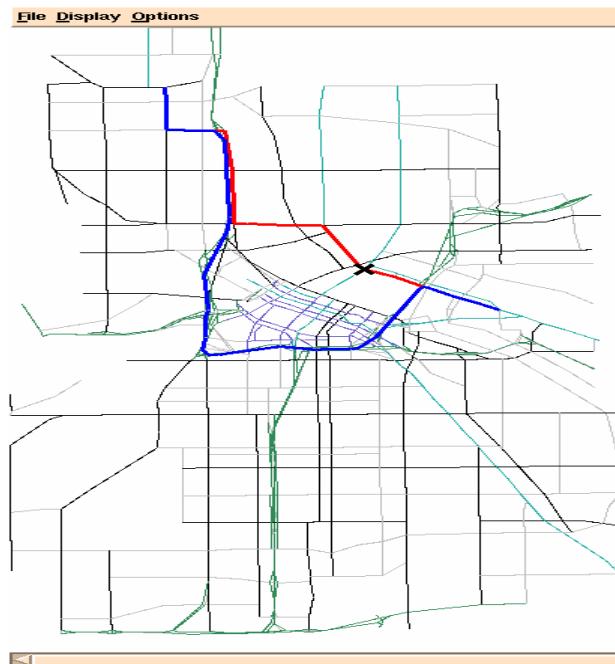
River Network & Queries

- List the names of all direct and indirect tributaries (支流) of Mississippi river
- List the direct tributaries of Colorado
- Which rivers could be affected if there is a spill in North Platte river

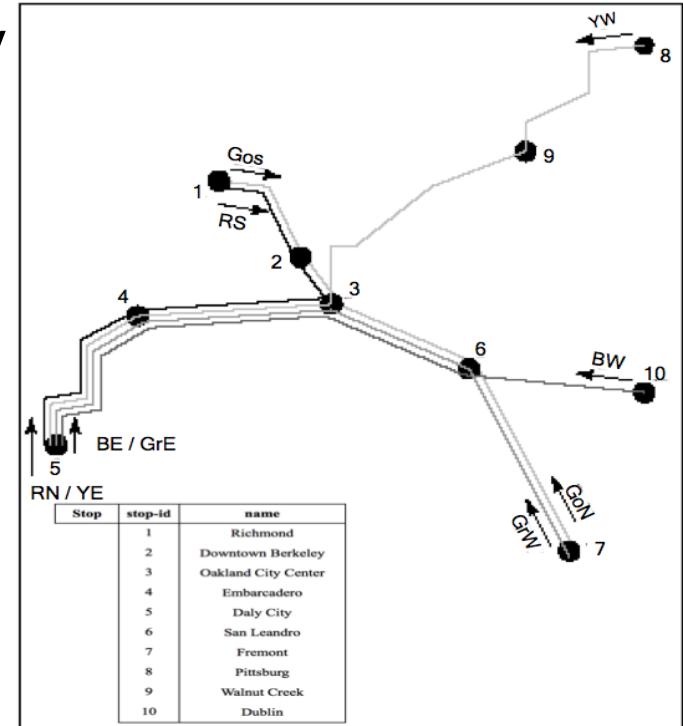


Spatial Networks: Three Examples

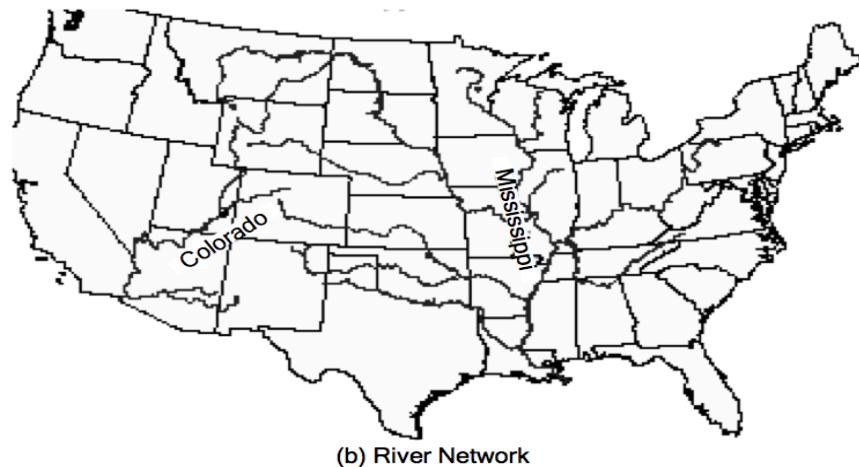
A Road Network



A Railway Network



A River Network



Quiz 3

- Which of the following are queries on a road network?
 - a) What is the shortest path to the nearest airport?
 - b) What is the driving distance to the nearest gas station?
 - c) How many road intersections are there in my city?
 - d) All of the above

第九章 空间网络模型与查询

- 9.1 Motivation and use cases
- 9.2 Conceptual model
- 9.3 Logical model
 - 9.3.1 Transitive Closure
 - 9.3.2 CONNECT statement
 - 9.3.3 RECURSIVE statement
 - 9.3.4 RECURSIVE in PostgreSQL
- 9.4 Physical model
 - 9.4.1 Storage and data structures
 - 9.4.2 Algorithms for connectivity query and shortest path
- 9.5 pgRouting

参考教材：

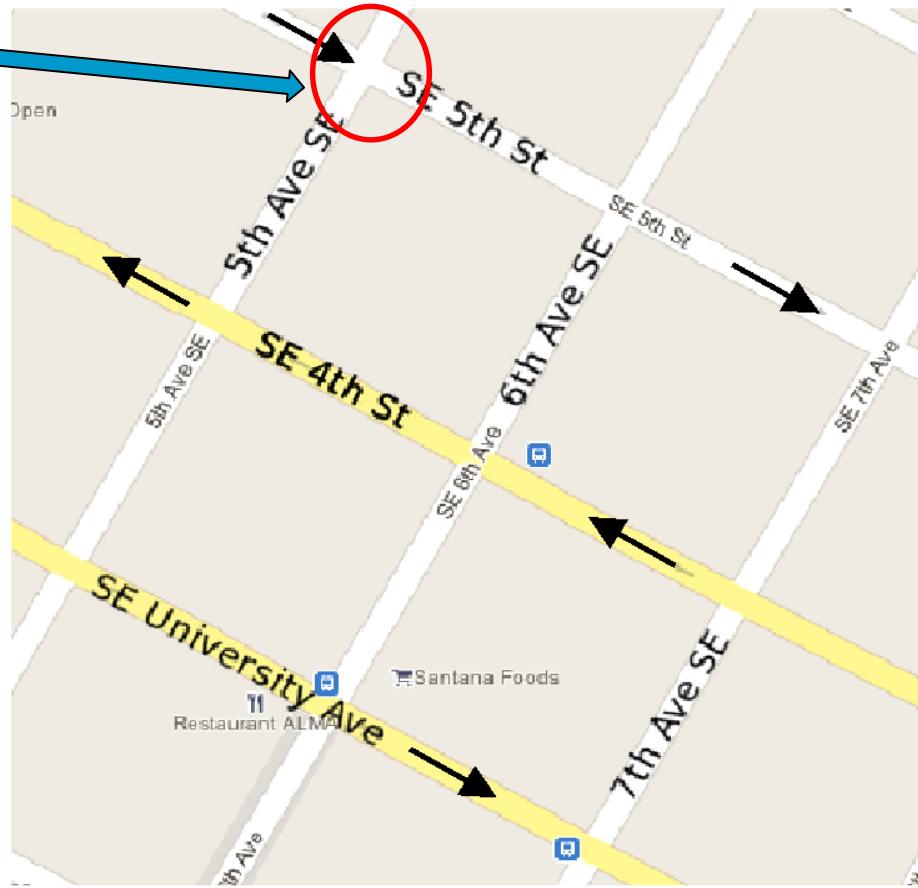
Spatial Databases: A Tour Chapter 6
空间数据库管理系统概论，3.3-3.4

Data Models of Spatial Networks

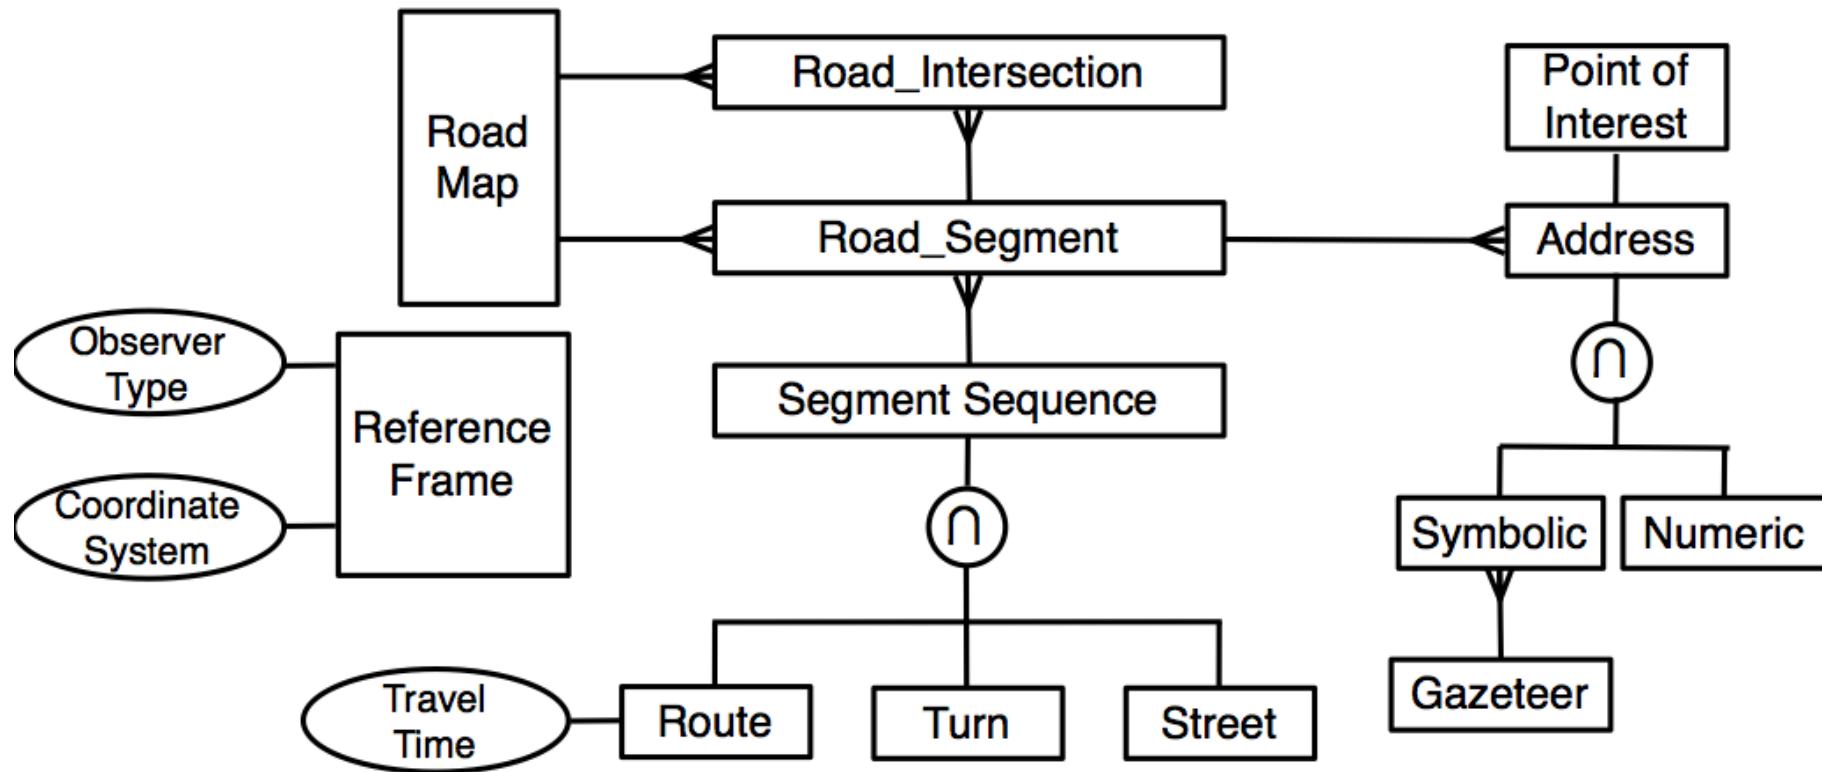
- Conceptual Model
 - Information Model: Entity Relationship Diagrams
 - Mathematical Model: Graphs
- Logical Data Model
 - Abstract Data types
 - Custom Statements in SQL
- Physical Data Model
 - Storage-Structures, File-Structures
 - Algorithms for common operations

Modeling Roadmaps

- Many Concepts, e.g.
 - Roads (or streets, avenues)
 - Road-Intersections
 - Road-Segments
 - Turns
 - ...

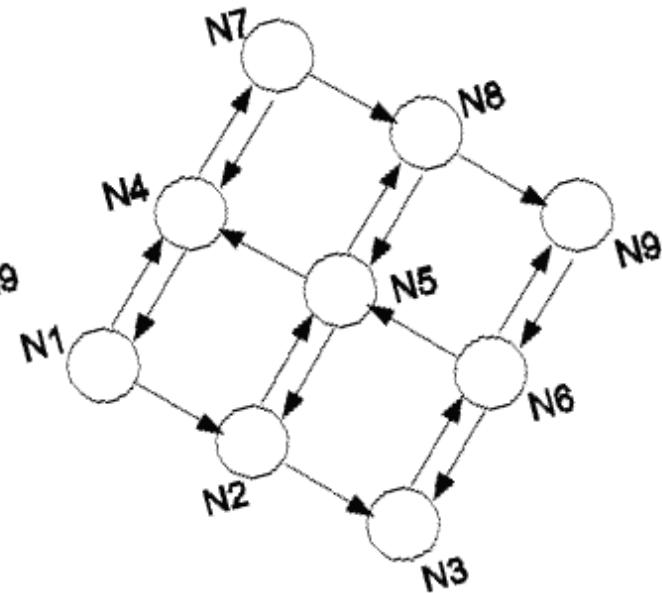
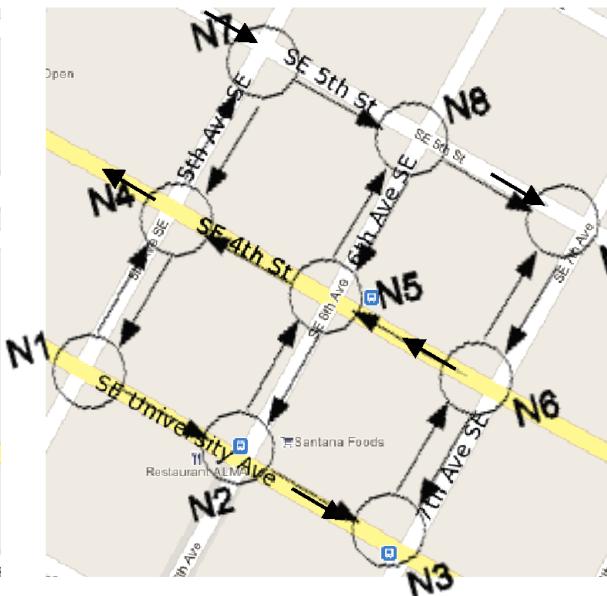
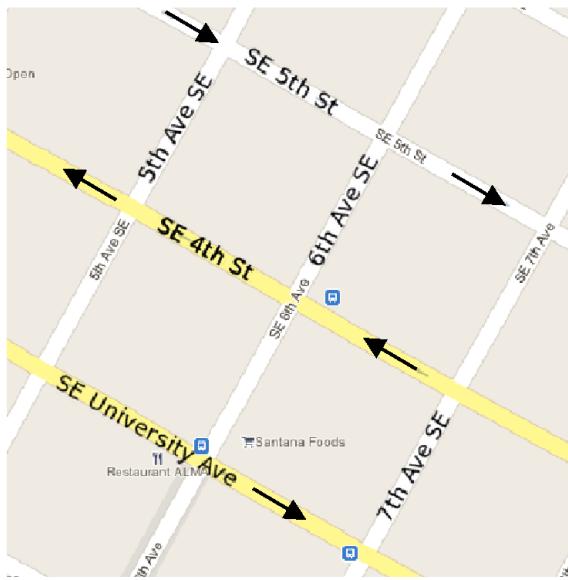


An Entity Relationship Diagram



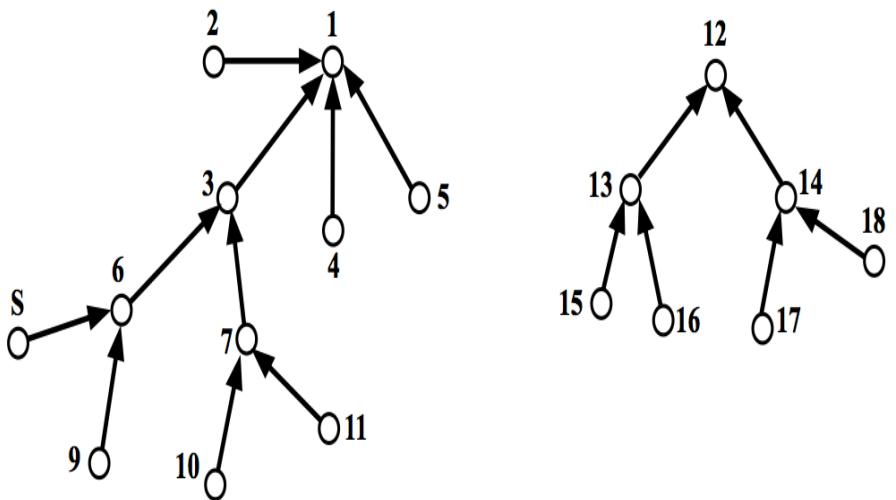
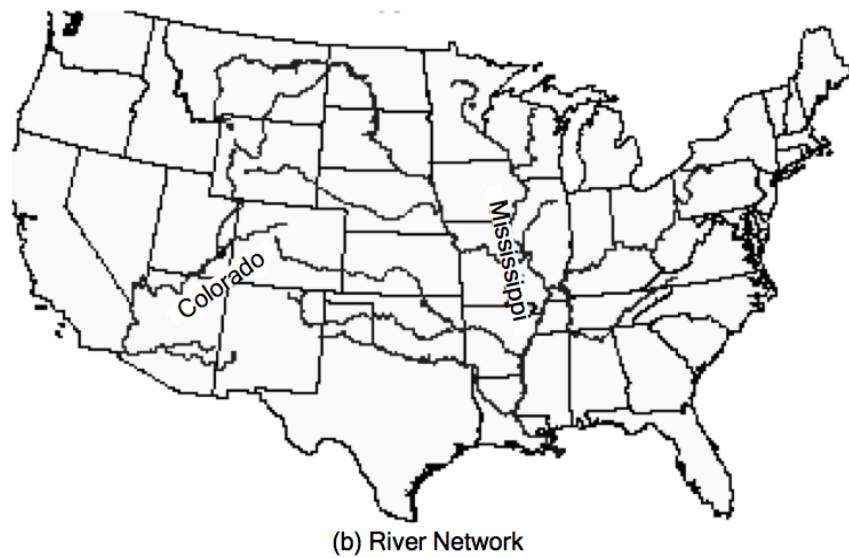
Graph Models

- A Simple Mathematical Model
 - A graph $G = (V, E)$
 - V = a finite set of vertices
 - E = a set of edges model a binary relationship between vertices
- Example



A Graph Model of River Network

- Nodes = rivers
- Edges = a river falls into another river



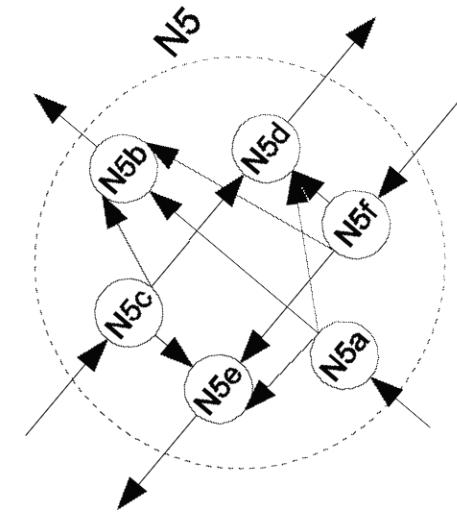
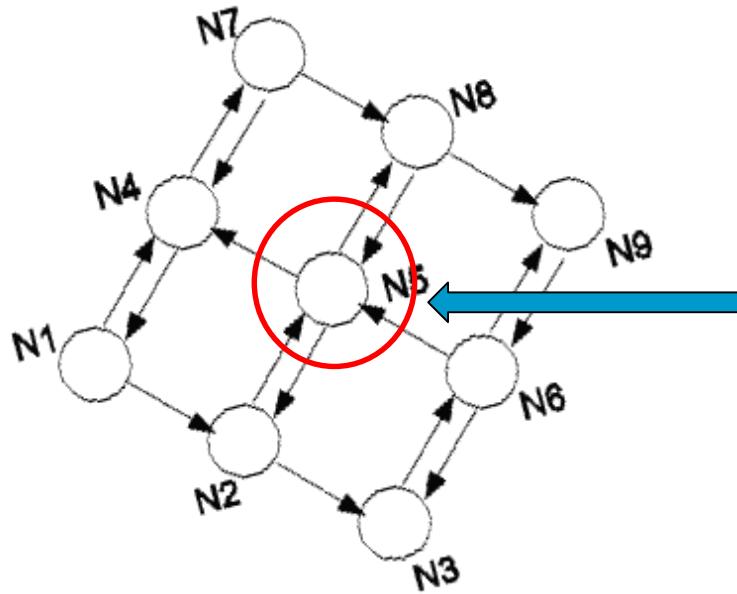
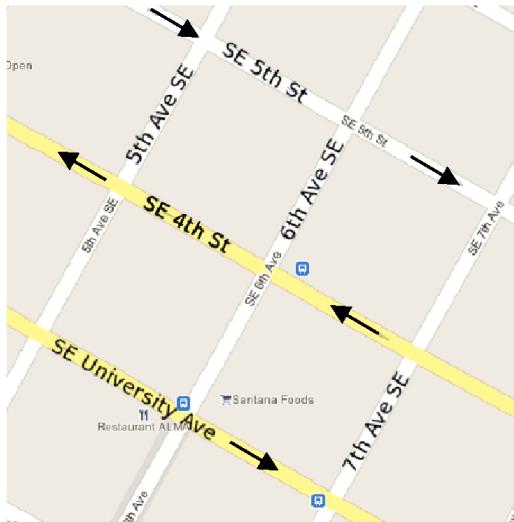
Pros and Cons of Graph Models

- Strength
 - Well developed mathematics for reasoning
 - Rich set of computational algorithms and data-structures
- Weakness
 - Models only one binary relationship
- Implications
 - A. Difficult to model multiple relationships, e.g., connect, turn
 - B. Multiple graph models possible for a spatial network



Modeling Turns in Roadmaps

- Approach 1: Model turns as a set of connects



- Approach 2: Use hyper-edges (and hyper-graphs)
- Approach 3: Annotate graph node with turn information

Alternative Graph Models for Roadmaps

- Choice 1:
 - Nodes = road-intersections
 - Edge (A, B) = road-segment **connects** adjacent road-intersections A, B
- Choice 2:
 - Nodes = (directed) road-segments
 - Edge (A, B) = **turn** from road-segment A to road-segment B
- Choice 3:
 - Nodes = roads
 - Edge(A, B) = road A **intersects_with** road B

Quiz 4

- Which of the following are usually not captured in common graph models of roadmaps?
 - a) Turn restrictions (e.g., no left turn)
 - b) Road intersections
 - c) Road segments
 - d) All of the above

第九章 空间网络模型与查询

- 9.1 Motivation and use cases
- 9.2 Conceptual model
- 9.3 Logical model
 - 9.3.1 Transitive Closure
 - 9.3.2 CONNECT statement
 - 9.3.3 RECURSIVE statement
 - 9.3.4 RECURSIVE in PostgreSQL
- 9.4 Physical model
 - 9.4.1 Storage and data structures
 - 9.4.2 Algorithms for connectivity query and shortest path
- 9.5 pgRouting

参考教材：

Spatial Databases: A Tour Chapter 6
空间数据库管理系统概论，3.3-3.4

Data Models of Spatial Networks

- Conceptual Model
 - Information Model: Entity Relationship Diagrams
 - Mathematical Model: Graphs
- **Logical Data Model & Query Languages**
 - Abstract Data types
 - Custom Statements in SQL
- Physical Data Model
 - Storage-Structures, File-Structures
 - Algorithms for common operations

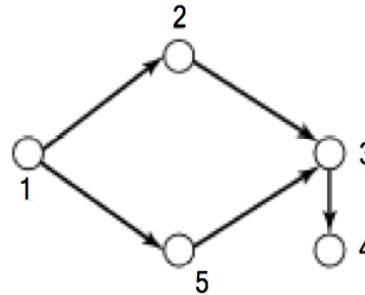
9.3.1 Transitive Closure

- Consider a graph $G = (V, E)$
- Transitive closure(G) = $G^* = (V^*, E^*)$, where
 - $V^* = V$
 - (A, B) in E^* if and only if there is a path from A to B in G

Transitive Closure - Example

- Example

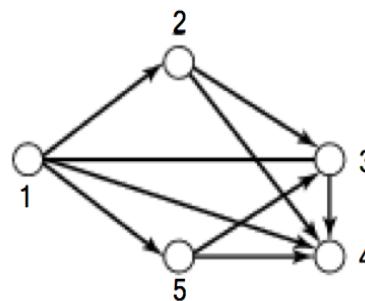
- G has 5 nodes and 5 edges
- G^* has 5 nodes and 9 edges
- Note edge (1,4) in G^* for
 - path (1, 2, 3, 4) in G



(a) Graph G

R	SOURCE	DEST
1	2	
1	5	
2	3	
3	4	
5	3	

(b) Relation form



(c) Transitive closure (G) = Graph G

X	SOURCE	DEST
1	2	
1	5	
2	3	
3	4	
5	3	
1	3	
2	4	
5	4	
1	4	

(d) Transitive closure in relation form

Limitations of Original SQL

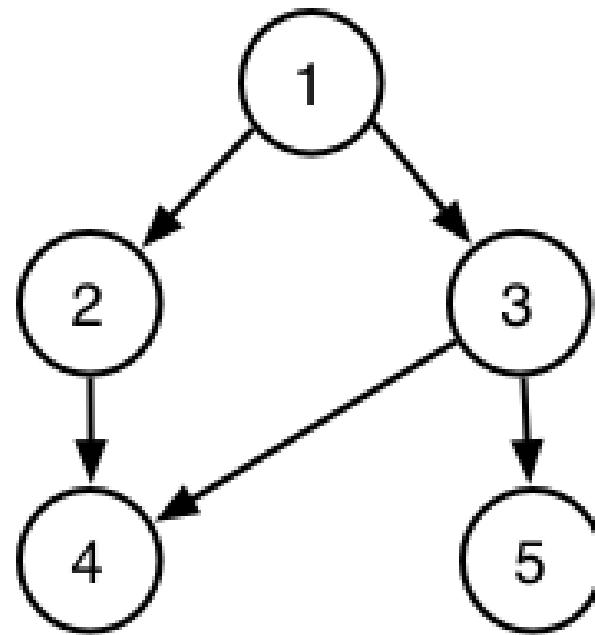
- Recall Relation Algebra based languages
 - Ex. original SQL
 - Can not compute transitive closure, e.g., shortest path

Supporting Graphs in SQL

- Abstract Data Type (user defined)
 - SQL3
 - May include shortest path operation!
- Custom Statements
 - SQL2 - CONNECT clause in SELECT statement
 - For **directed acyclic graphs**, e.g. hierarchies
 - SQL3 - WITH RECURSIVE statement
 - Transitive closure on general graphs
 - SQL3 - User defined data types
 - Can include shortest path operation!

Quiz 5

- Which of the following is not in the transitive closure of the following graph?
 - a) (1, 5)
 - b) (1, 4)
 - c) (2, 3)
 - d) (3, 4)

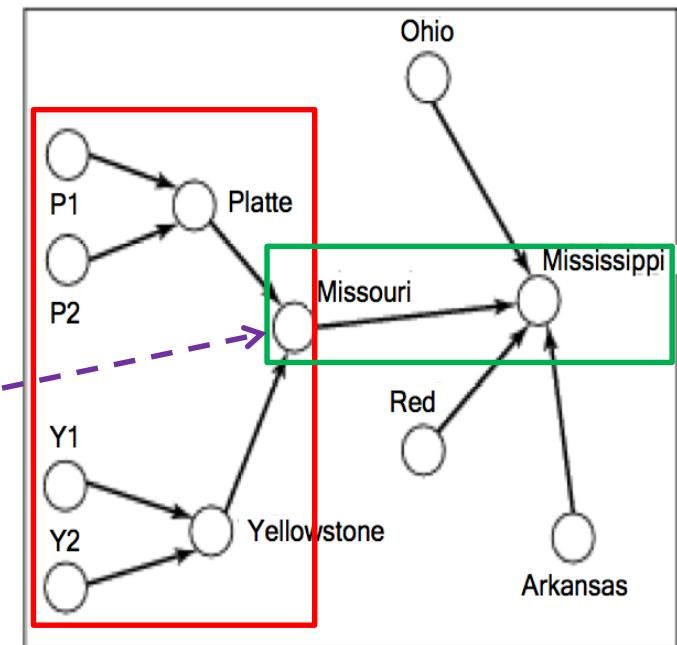


9.3.2 Querying Graphs: Overview

- Relational Algebra
 - Can not express transitive closure queries
- Two ways to extend SQL to support graphs
 - Abstract Data Types
 - Custom Statements
 - SQL2 - CONNECT BY clause(s) in SELECT statement
 - SQL3 - WITH RECURSIVE statement
 - SQL3 - User defined data types

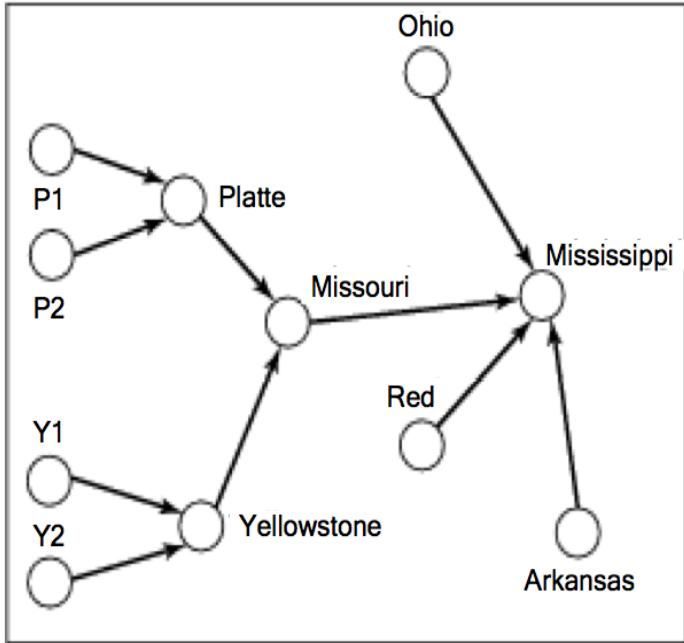
CONNECT BY: Input, Output

- Input:
 - (a) Edges of a directed acyclic graph G
 - (b) Start Node S, e.g., Missouri
 - (c) Travel Direction
- Output: Transitive closure of G
 - Ex. Predecessors of S = Missouri
 - Ex. Successors of S = Missouri
(密苏里)



(a) Mississippi network (Y1 = Bighorn river,
Y2 = Power river, P1 = Sweet water River,
P2 = Big Thompson river)

Directed Edges: Tabular Representation



(a) Mississippi network (Y1 = Bighorn river,
Y2 = Power river, P1 = Sweet water River,
P2 = Big Thompson river)

Table: Falls_Into

Source	Dest
P1	Platte
P2	Platte
Y1	Yellowstone
Y2	Yellowstone
Platte	Missouri
Yellowstone	Missouri
Missouri	Mississippi
Ohio	Mississippi
Red	Mississippi
Arkansas	Mississippi

CONNECT BY - PRIOR - START WITH

```
SELECT source  
FROM Falls_Into ←  
CONNECT BY PRIOR source = dest  
START WITH dest = "Missouri"
```

Table: Falls_Into	
Source	Dest
P1	Platte
P2	Platte
Y1	Yellowstone
Y2	Yellowstone
Platte	Missouri
Yellowstone	Missouri
Missouri	Mississippi
Ohio	Mississippi
Red	Mississippi
Arkansas	Mississippi

- Q? What does CONNECT BY ...
PRIOR specify?
- Direction of travel
 - Example: **From Dest to Source**
 - Alternative: **From Source to Dest**

CONNECT BY- PRIOR - START WITH

- Syntax details
 - FROM clause a table for directed edges of an acyclic graph
 - PRIOR identifies direction of traversal for the edge
 - START WITH specifies first vertex for path computations
- Semantics
 - List all nodes reachable from first vertex using directed edge in specified table
 - Assumption - no cycle in the graph!
 - Not suitable for train networks, road networks

CONNECT BY - PRIOR - START WITH

Choice 1: Travel from Dest to Source

Ex. List direct & indirect tributaries of Missouri.

```
SELECT source  
FROM Falls_Into  
CONNECT BY PRIOR source = dest  
START WITH dest = "Missouri"
```

Choice 2: Travel from Source to Dest

Ex. Which rivers are affected by spill in Missouri?

```
SELECT dest  
FROM Falls_Into  
CONNECT BY source = PRIOR dest  
START WITH source = "Missouri"
```

Table: Falls_Into

Source	Dest
P1	Platte
P2	Platte
Y1	Yellowstone
Y2	Yellowstone
Platte	Missouri
Yellowstone	Missouri
Missouri	Mississippi
Ohio	Mississippi
Red	Mississippi
Arkansas	Mississippi

Execution Trace – Step 1

```
SELECT source  
FROM Falls_Into  
CONNECT BY PRIOR source = dest  
START WITH dest = "Missouri"
```

1. Prior Result = SELECT * FROM Falls_Into WHERE dest = "Missouri"

Table: "Prior "	
Source	Dest
Platte	Missouri
Yellowstone	Missouri

Table: Falls_Into

Source	Dest
P1	Platte
P2	Platte
Y1	Yellowstone
Y2	Yellowstone
Platte	Missouri
Yellowstone	Missouri
Missouri	Mississippi
Ohio	Mississippi
Red	Mississippi
Arkansas	Mississippi

Execution Trace – Step 2

```
SELECT source  
FROM Falls_Into  
CONNECT BY PRIOR source = dest  
START WITH dest = "Missouri"
```

2. Iteratively add `Join(Prior_Result.source = Falls_Into.dest)`

Prior Result = `SELECT * FROM Falls_Into
WHERE Prior_Result.source = Falls_Into.dest`

Prior Result

Source
P1
P2
Y1
Y2
Platte
Yellowstone

Prior Result	
Source	Dest
Platte	Missouri
Yellowstone	Missouri

Table: Falls_Into

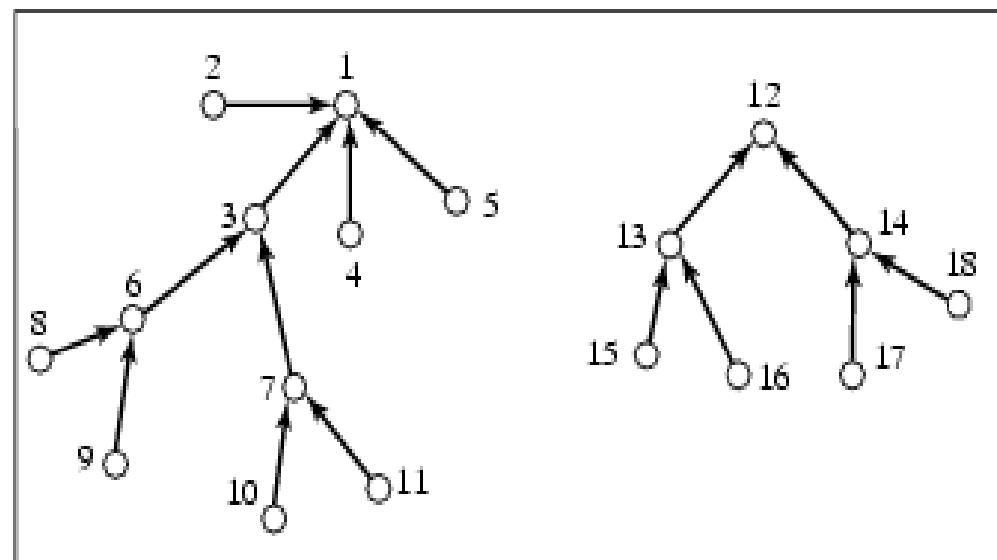
Source	Dest
P1	Platte
P2	Platte
Y1	Yellowstone
Y2	Yellowstone
Platte	Missouri
Yellowstone	Missouri
Missouri	Mississippi
Ohio	Mississippi
Red	Mississippi
Arkansas	Mississippi

SQL CONNECT Exercise

- Study 2 SQL queries on right
 - Note different use of PRIOR keyword
- Compute results of each query
- Which one returns ancestors of 3?
- Which returns descendants of 3?
- Which query lists river affected by oil spill in Missouri (id = 3)?

```
SELECT source FROM FallsInto  
CONNECT BY PRIOR source = dest  
START WITH dest = 3
```

```
SELECT dest FROM FallsInto  
CONNECT BY source = PRIOR dest  
START WITH source = 3
```



Quiz 6

- Which of the following is false about CONNECT BY clause?
 - a) It is only able to output predecessors, but not successors, of the start node
 - b) It is able to output transitive closure of a directed graph
 - c) It usually works with PRIOR and START WITH keywords
 - d) None of the above

9.3.3 Querying Graphs: Overview

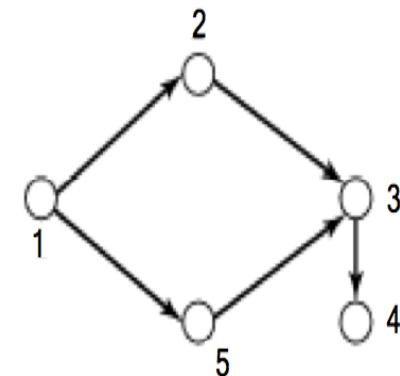
- Relational Algebra
 - Can not express transitive closure queries
- Two ways to extend SQL to support graphs
 - Abstract Data Types
 - Custom Statements
 - SQL2 - CONNECT BY clause(s) in SELECT statement
 - **SQL3 - WITH RECURSIVE statement**
 - SQL3 - User defined data types

WITH RECURSIVE: Input, Output

- Input:
 - (a) Edges of a directed graph G
 - (b) Sub-queries to
 - Initialize results
 - Recursively grow results
 - Additional constraints

R	SOURCE	DEST
	1	2
	1	5
	2	3
	3	4
	5	3

(b) Relation form



(a) Graph G

- Output: Transitive closure of G
 - Ex. Predecessors of a node
 - Ex. Successors of a node

Syntax of WITH RECURSIVE Statement

WITH RECURSIVE X(source, dest) Description of Result Table
AS (SELECT source, dest FROM R) Initialization Query

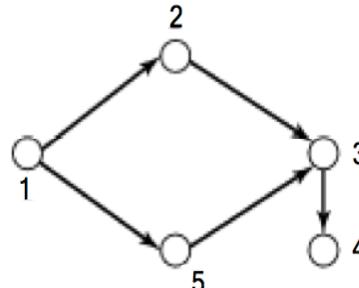
UNION

(SELECT R.source, X.dest
FROM R, X
WHERE R.dest = X.source)

← Recursive Query to grow result

Example Input and Output

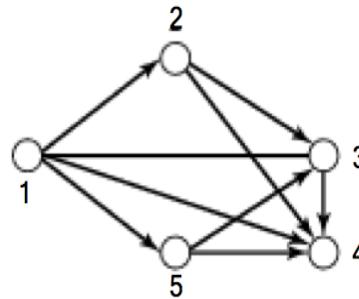
WITH RECURSIVE X(source,dest)
AS (SELECT source,dest FROM R)
UNION
(SELECT R.source, X.dest
FROM R, X
WHERE R.dest=X.source)



(a) Graph G

R	SOURCE	DEST
1	2	
1	5	
2	3	
3	4	
5	3	

(b) Relation form



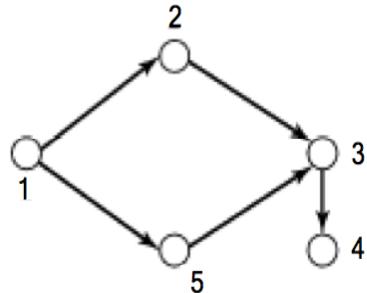
(c) Transitive closure (G) = Graph G

X	SOURCE	DEST
1	2	
1	5	
2	3	
3	4	
5	3	
1	3	
2	4	
5	4	
1	4	

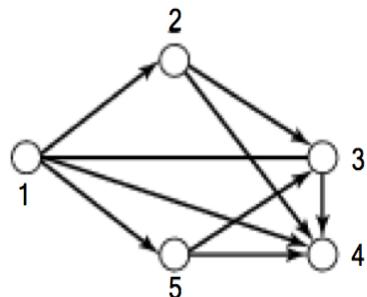
(d) Transitive closure in relation form

SQL3 Recursion Example -

Meaning



(a) Graph G



(c) Transitive closure (G) = Graph G

R	
SOURCE	DEST
1	2
1	5
2	3
3	4
5	3

(b) Relation form

X	
SOURCE	DEST
1	2
1	5
2	3
3	4
5	3
1	3
2	4
5	4
1	4

(d) Transitive closure in relation form

- Initialize X by
(SELECT source,dest FROM R)
- Recursively grow X by
(SELECT R.source, X.dest
FROM R, X
WHERE R.dest=X.source)
- Infer X(a,c) from R(a,b),X(b,c)
 - Infer X(1,3) from R(1,2),X(2,3)
 - Infer X(2,4) from R(2,3),X(3,4)
 - Infer X(5,4) from R(5,3),X(3,4)
 - Infer X(1,4) from R(1,5),X(5,4)

SQL3 Recursion

- Syntax
 - WITH RECURSIVE <Relational Schema>
 - AS <Query to populate relational schema> Syntax details
 - <Relational Schema> lists columns in result table with directed edges
 - <Query to populate relational schema> has UNION of nested sub-queries
 - Base cases to initialize result table
 - Recursive cases to expand result table

SQL3 Recursion

- Semantics
 - Results relational schema say $X(\text{source}, \text{dest})$
 - Columns source and dest come from same domain, e.g. Vertices
 - X is a edge table, $X(a,b)$ directed from a to b
 - Result table X is initialized using base case queries
 - Result expanded using $X(a, b)$ and $X(b, c)$ implies $X(a, c)$

With Recursive for Connectivity

- Connect by
 - For directed acyclic graphs, e.g. hierarchies
 - PostgreSQL, SQL Server not supported
 - Oracle supported
- With Recursive
 - Transitive closure on general graphs
 - PostgreSQL, SQL Server, Oracle supported

<http://www.postgresql.org/docs/current/static/queries-with.html>

Case Studies

- Goal: Compare relational schemas for spatial networks
 - River networks has an edge table, Falls_Into
 - BART train network does not have an edge table
 - Edge table is crucial for using SQL transitive closure
 - Exercise: Proposed a different set of table to model BART as a graph
 - using an edge table connecting stops
- River networks - graph model
 - Can use SQL transitive closure to compute ancestors or descendants of a river
 - We saw examples using CONNECT BY clause
 - Exercises explore use of WITH RECURSIVE statement

Case Studies

- BART train network - non-graph model
 - Entities = Stop, Route
 - Relationship = aMemberOf(Stop, Route)
 - Can not use SQL recursion
 - No table can be viewed as edge table
 - RouteStop table is a subset of transitive closure
- Transitive closure queries on edge(from_stop, to_stop)
 - A few can be answered by querying RouteStop table
 - Many can not be answered
 - Find all stops reachable from Downtown Berkeley

Quiz 7

- Which of the following are true about WITH RECURSIVE clause?
 - a) It is able to output transitive closure of a directed graph
 - b) It usually works with an edge table
 - c) It includes two SELECT statements
 - d) All of the above

9.3.4 With Recursive in PostgreSQL

- With clause
 - A way to write auxiliary statements for use in a larger query
 - Define temporary tables that exist just for one query
- A WITH query can refer to its own output
- Sum the integers from 1 through 100

WITH RECURSIVE t(n) AS (

VALUES (1)

UNION ALL

SELECT n+1 FROM t WHERE n < 100

)

SELECT sum(n) FROM t;

A non-recursive term

Union or Union ALL

A recursive term

Recursive Query Evaluation

- Step 1: Evaluate the non-recursive term
 - For UNION, discard duplicate rows
 - Include all remaining rows in **the result of the recursive query**, and also place them in **a temporary working table**
 - Example
 - VALUES (1)

```
WITH RECURSIVE t(n) AS (
    VALUES (1)
    UNION ALL
    SELECT n+1 FROM t WHERE n < 100)
```

Query Result	Working
1	1

Recursive Query Evaluation

- Step 2: So long as working table is not empty, repeat these steps:
 - 2.1 Evaluate the recursive term, substituting the current contents of **the working table** for the recursive self-reference
 - Example
 - SELECT n+1 FROM **t** WHERE n < 100
 - t is the query result table Q? or the working table W

```
WITH RECURSIVE t(n) AS (
    VALUES (1)
    UNION ALL
    SELECT n+1 FROM t WHERE n < 100)
```

Intermediate	Query Q	Working W
2	1	1

Recursive Query Evaluation

- Step 2: So long as working table is not empty, repeat these steps:
 - 2.1 Evaluate the recursive term, substituting the current contents of **the working table** for the recursive self-reference
 - 2.2 For UNION, discard duplicated rows and rows that duplicated any previous result row
 - Where are previous result rows?

```
WITH RECURSIVE t(n) AS (
    VALUES (1)
    UNION ALL
    SELECT n+1 FROM t WHERE n < 100)
```

Intermediate	Query Q	Working W
2	1	1

Recursive Query Evaluation

- Step 2: So long as working table is not empty, repeat these steps:
 - 2.1 Evaluate the recursive term, substituting the current contents of **the working table** for the recursive self-reference
 - 2.2 For UNION, discard duplicated rows and rows that duplicated any previous result row
 - 2.3 Including all remaining rows in **the result of the recursive query**, and also place them in **a temporary intermediate table**

Intermediate	Query Q	Working W
2	1	1
	2	

Recursive Query Evaluation

- Step 2: So long as working table is not empty, repeat these steps:
 - 2.1 Evaluate the recursive term, substituting the current contents of **the working table** for the recursive self-reference
 - 2.2 For UNION, discard duplicated rows and rows that duplicated any previous result row
 - 2.3 Including all remaining rows in **the result of the recursive query**, and also place them in **a temporary intermediate table**
 - 2.4 Replace the contents of **the working table** with the contents of **the intermediate table**, then empty **the intermediate table**

Recursive Query Evaluation

- Step 2: So long as working table is not empty, repeat these steps
 - $\text{SELECT } n+1 \text{ FROM } W \text{ WHERE } n < 100$

Intermediate	Query Q	Working W
	1	2
	2	
Intermediate	Query Q	Working W
3	1	2
	2	
Intermediate	Query Q	Working W
	1	3
	2	
	3	

Recursive Query Evaluation

- Step 2: So long as working table is not empty, repeat these steps
 - $\text{SELECT } n+1 \text{ FROM } W \text{ WHERE } n < 100$

```
WITH RECURSIVE t(n) AS (
    VALUES (1)
    UNION ALL
    SELECT n+1 FROM t WHERE n < 100)
```

t
1
.....
100

Intermediate	Query Q	Working W
	1	100
	
	100	

Intermediate	Query Q	Working W
	1	
	
	100	

Recursive Query Evaluation in C

- with recursive $X(\dots)$ as (
 $\text{SQL_A}(O)$
 $\text{union} / \text{union all}$
 $\text{SQL_B}(O, X))$
 - $X(\dots) = \text{SQL_A}[\text{remove duplicates for union};$
 $W(\dots) = X(\dots);$
 $\text{while (table } W \text{ is not empty) } \{$
 $T(\dots) = \text{SQL_B}(O, W) [\text{except } X(\dots) \text{ for union}] ;$
 $X(\dots) = X(\dots) \text{ union / union all } T(\dots);$
 $W(\dots) = T(\dots);$
- }
- O: 原始关系/表
X: 最终关系/表
W: 临时关系/表
T: 临时关系/表

问题: Union与Union ALL区别
避免使用Union ALL

Recursive Query Evaluation in C

WITH RECURSIVE t(n) AS (

VALUES (1)

UNION ALL

SELECT n+1 FROM t WHERE n < 100)

- insert into t(n) values(1);
delete from W; insert into W(n) select * from t;
while ((select count(*) from W) <> 0) {
 delete from T;
 insert into T(n) select n+1 from W where n < 100;
 insert into t(n) select * from T; // union all
 delete from W; insert into W from select * from T;
}

t: 最终关系/表

W: 临时关系/表

T: 临时关系/表

With Recursive Limitation

- Strictly speaking, this process is **iteration** not recursion, but **RECURSIVE** is the terminology chosen by the SQL standards committee
- Recursive queries are typically used to deal with **hierarchical** or **tree-structured data**

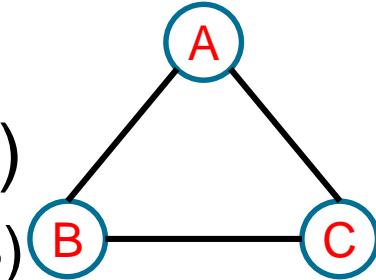
With Recursive Limitation

- Important: the recursive part of the query will eventually return no tuples, or else the query will **loop indefinitely**
 - Using UNION instead of UNION ALL can accomplish this by discarding rows that duplicate previous output rows
- A cycle does not involve output rows that are completely duplicate
 - It may be necessary to check just one or a few fields to see if the same point has been reached before
- Standard method
 - Compute **an array** of the already-visited values

With Recursive Example 1

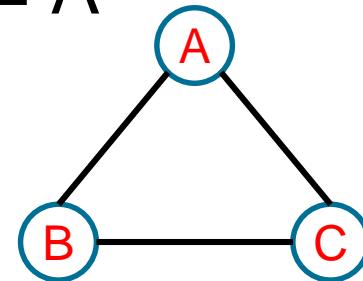
- 广度优先遍历 - 深度depth
- 数据库中，右图的关系为edges(start, end)
 - 6行记录(A,B), (A,C), (B,A), (B,C), (C,A), (C,B)
- with recursive X(node, depth) as (

```
select start, 0 from edges where start = A
union
select end, depth + 1 from edges, X
where start = node and depth < 3)
```
- 如果不加depth<3， 上述语句会死循环



With Recursive Example 1

- with recursive $X(\text{node}, \text{depth})$ as (
 select start, 0 from edges where start = A
 union
 select end, depth + 1 from edges, X
 where start = node and depth < 3)



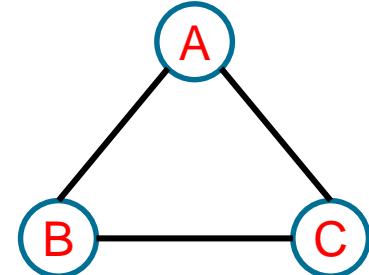
- 初始 X : $(A, 0)$
- 第1次迭代: $(A, 0), (B, 1), (C, 1)$
- 第2次迭代: $(A, 0), (B, 1), (C, 1), (A, 2), (B, 2), (C, 2)$
 - 如果使用union all得到: $(A, 2), (C, 2), (A, 2), (B, 2)$
- 第3次迭代: $(A, 0), (B, 1), (C, 1), (A, 2), (B, 2), (C, 2)$
 $(A, 3), (B, 3), (C, 3)$
-

With Recursive Example 2

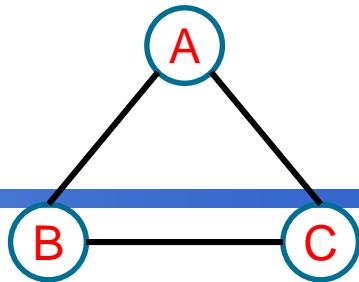
- 避免重复遍历节点path和cycle
- with recursive X(node, depth, path, cycle) as (
 select start, 0, array[start], false from edges where start = A
 union
 select end, depth + 1, path || end, end = any(path)
 from edges, X
 where start = node and not cycle)
- PostgreSQL 中的数组
 - 初始化: array[...]
 - 增加元素: path || end
 - 判断元素是否在数组中: end = any(path)

With Recursive Example 2

- with recursive `X(node, depth, path, cycle)` as (
 `select start, 0, array[start], false from edges where start = A`
 `union`
 `select end, depth + 1, path || end, end = any(path)`
 `from edges, X`
 `where start = node and not cycle`)
- 初始X: `(A, 0, [A], false)`
- 第1次迭代: `(A, 0, [A], false)`, `(B, 1, [A, B], false)`, `(C, 1, [A, C], false)`
- 第2次迭代: `(A, 0, [A], false)`, `(B, 1, [A, B], false)`, `(C, 1, [A, C], false)`,
`(A, 2, [A, B, A], true)`, `(C, 2, [A, B, C], false)`,
`(A, 2, [A, C, A], true)`, `(B, 2, [A, C, B], false)`
- 第3次迭代: `(A, 0, [A], false)`, `(B, 1, [A, B], false)`, `(C, 1, [A, C], false)`,
`(A, 2, [A, B, A], true)`, `(C, 2, [A, B, C], false)`,
`(A, 2, [A, C, A], true)`, `(B, 2, [A, C, B], false)`,
`(A, 3, [A, B, C, A], true)`, `(B, 3, [A, B, C, B], true)`,
`(A, 3, [A, C, B, A], true)`, `(C, 3, [A, C, B, C], true)`



With Recursive Example 2



- 初始X: (A, 0, [A], false)
- 第1次迭代: (A, 0, [A], false), (B, 1, [A, B], false), (C, 1, [A, C], false)
- 第2次迭代: (A, 0, [A], false), (B, 1, [A, B], false), (C, 1, [A, C], false),
(A, 2, [A, B, A], true), (C, 2, [A, B, C], false),
(A, 2, [A, C, A], true), (B, 2, [A, C, B], false)
- 第3次迭代: (A, 0, [A], false), (B, 1, [A, B], false), (C, 1, [A, C], false),
(A, 2, [A, B, A], true), (C, 2, [A, B, C], false),
(A, 2, [A, C, A], true), (B, 2, [A, C, B], false),
(A, 3, [A, B, C, A], true), (B, 3, [A, B, C, B], true),
(A, 3, [A, C, B, A], true), (C, 3, [A, C, B, C], true)
- 第4次迭代: (A, 0, [A], false), (B, 1, [A, B], false), (C, 1, [A, C], false),
(A, 2, [A, B, A], true), (C, 2, [A, B, C], false),
(A, 2, [A, C, A], true), (B, 2, [A, C, B], false),
(A, 3, [A, B, C, A], true), (B, 3, [A, B, C, B], true),
(A, 3, [A, C, B, A], true), (C, 3, [A, C, B, C], true)
- 结束

With Recursive Example 3

- Find all the direct and indirect sub-parts of a product, given only a table that shows immediate inclusions
 - parts(part, sub_part, quantity)
 - Similar to Rivers

With Recursive Example 3

- Find all the direct and indirect sub-parts of a product

```
WITH RECURSIVE included_parts(sub_part, part, quantity) AS (
  SELECT sub_part, part, quantity FROM parts WHERE part = 'our_product'
  UNION ALL
  SELECT p.sub_part, p.part, p.quantity FROM included_parts pr, parts p
  WHERE p.part = pr.sub_part
)
```

```
SELECT sub_part, SUM(quantity) as total_quantity
FROM included_parts
GROUP BY sub_part
```

With Recursive Example 4

- Graph search: graph(id, link, data)

```
WITH RECURSIVE search_graph(id, link, data, depth) AS (
```

```
    SELECT g.id, g.link, g.data, 1
```

```
    FROM graph g
```

```
    UNION ALL
```

```
    SELECT sg.id, sg.link, sg.data, sg.depth + 1
```

```
    FROM graph g, search_graph sg
```

```
    WHERE g.id = sg.link
```

```
)
```

```
SELECT * FROM search_graph;
```

- This query will loop if the link relationship contain cycle

- UNION ALL → UNION

- Would not eliminate the looping due to the “depth” output

With Recursive Example 4

- Need to recognize whether we have reached the same row again while following path of links
 - Add two columns path and cycle to the loop-prone query

```
WITH RECURSIVE search_graph(id, link, data, depth, path, cycle) AS (
```

```
    SELECT g.id, g.link, g.data, 1,  
          ARRAY[g.id],  
          false
```

Represent the “path” taken to
reach any particular row

```
    FROM graph g
```

```
    UNION ALL
```

```
    SELECT sg.id, g.link, g.data, sg.depth + 1,  
          path || g.id,  
          g.id = ANY(path)
```

```
    FROM graph g, search_graph sg
```

```
    WHERE g.id = sg.link AND NOT cycle
```

```
)
```

With Recursive Example 4

- Generally, more than one field needs to be checked to recognize a cycle, use an array of rows

```
WITH RECURSIVE search_graph(id, link, data, depth, path, cycle) AS (
    SELECT g.id, g.link, g.data, 1,
           ARRAY[ROW(g.f1, g.f2)],
           false
      FROM graph g
 UNION ALL
    SELECT g.id, g.link, g.data, sg.depth + 1,
           path || ROW(g.f1, g.f2),
           ROW(g.f1, g.f2) = ANY(path)
      FROM graph g, search_graph sg
     WHERE g.id = sg.link AND NOT cycle
)
SELECT * FROM search_graph;
```

Recursive in PostgreSQL

- A helpful trick for testing queries when you are not certain if they might loop is to place a **LIMIT** in the parent query

```
WITH RECURSIVE t(n) AS (
    VALUES (1)
    UNION ALL
    SELECT n+1 FROM t WHERE n < 100
)
SELECT sum(n) FROM t LIMIT 100;
```

- PostgreSQL evaluates only as many rows of a WITH query as are actually fetched by the parent query
 - Other DBMS might work different
 - Won't work if order by or join them to some other table

避免死循环方法

1. Union
2. Limit
3. Depth控制
4. Cycle检测

With Queries Applications

- Avoid redundant work
 - Evaluated only once per execution of the parent query, even if they are referred to more than once by the parent query or sibling WITH queries
 - Expensive calculations that are needed in multiple places can be placed within a WITH query
- Prevent unwanted multiple evaluations of functions with side-effects
 - The optimizer is less able to push restrictions from the parent query down into a WITH query than an ordinary sub-query

第九章 空间网络模型与查询

- 9.1 Motivation and use cases
- 9.2 Conceptual model
- 9.3 Logical model
 - 9.3.1 Transitive Closure
 - 9.3.2 CONNECT statement
 - 9.3.3 RECURSIVE statement
 - 9.3.4 RECURSIVE in PostgreSQL
- 9.4 Physical model
 - 9.4.1 Storage and data structures
 - 9.4.2 Algorithms for connectivity query and shortest path
- 9.5 pgRouting

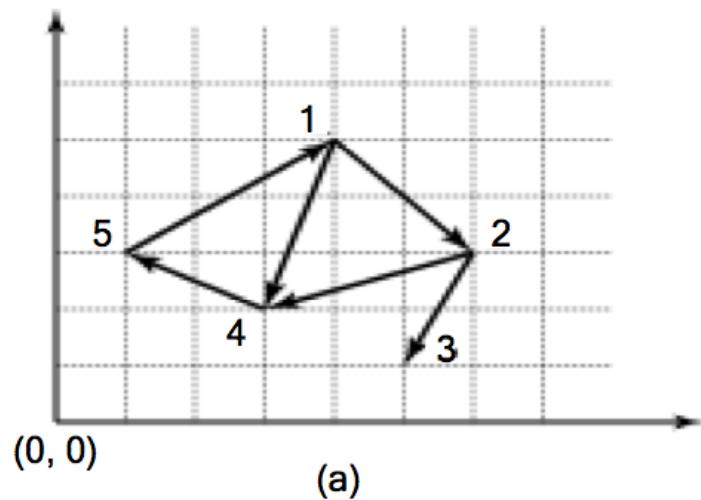
参考教材：
Spatial Databases: A Tour Chapter 6
空间数据库管理系统概论，3.3-3.4

Data Models of Spatial Networks

- Conceptual Model
 - Entity Relationship Diagrams, Graphs
- Logical Data Model & Query Languages
 - Abstract Data types
 - Custom Statements in SQL
- Physical Data Model
 - Storage: Data-Structures, File-Structures
 - Algorithms for common operations

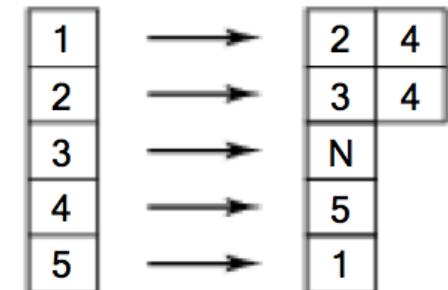
9.4.1 Main Memory Data-Structures

- Adjacency matrix
 - $M[A, B] = 1$ if and only if edge(vertex A, vertex B) exists
- Adjacency list
 - Maps a vertex to a list of its successors



		Destination				
		1	2	3	4	5
source	1	0	1	0	1	0
	2	0	0	1	1	0
	3	0	0	0	0	0
	4	0	0	0	0	1
	5	1	0	0	0	0

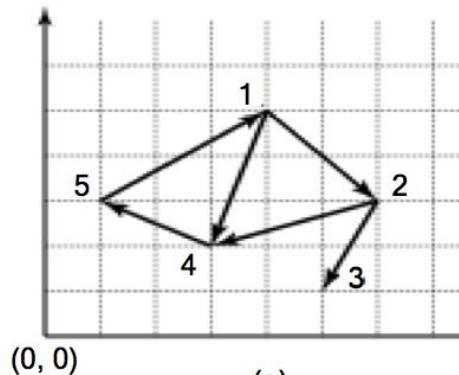
(b) Adjacency-Matrix



(c) Adjacency-List

Disk-based Tables

- Normalized tables
 - One for vertices, other for edges
- Denormalized
 - One table for nodes with adjacency lists



Node (R)		
id	x	y
1	4.0	5.0
2	6.0	3.0
3	5.0	1.0
4	3.0	2.0
5	1.0	3.0

Edge (S)		
source	dest	distance
1	2	$\sqrt{8}$
1	4	$\sqrt{10}$
2	3	$\sqrt{5}$
2	4	$\sqrt{10}$
4	5	$\sqrt{5}$
5	1	$\sqrt{18}$

(d) Node and Edge Relations

id	x	y	Successors	Predecessors
1	4.0	5.0	(2,4)	(5)
2	6.0	3.0	(3,4)	(1)
3	5.0	1.0	()	(2)
4	3.0	2.0	(5)	(1,2)
5	1.0	3.0	(1)	(4)

(e) Denormalized Node Table

Spatial Network Storage

- Problem Statement
 - Given a spatial network
 - Find efficient data-structure to store it on disk sectors
 - Goal - Minimize I/O-cost of operations
 - Find(), Insert(), Delete(), Create()
 - Get-A-Successor(), Get-Successors()
 - Constraints
 - Spatial networks are much larger than main memories
- Problems with geometric indices, e.g. R-tree
 - Clusters objects by proximity not edge connectivity
 - Performs poorly if edge connectivity not correlated with proximity (邻近)
- Trends: graph based methods

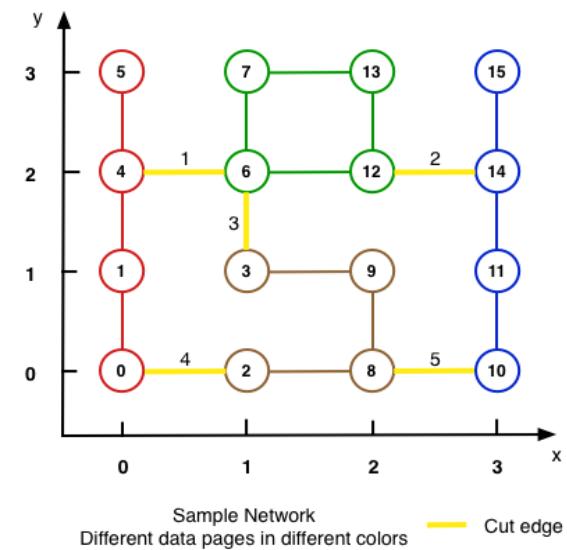
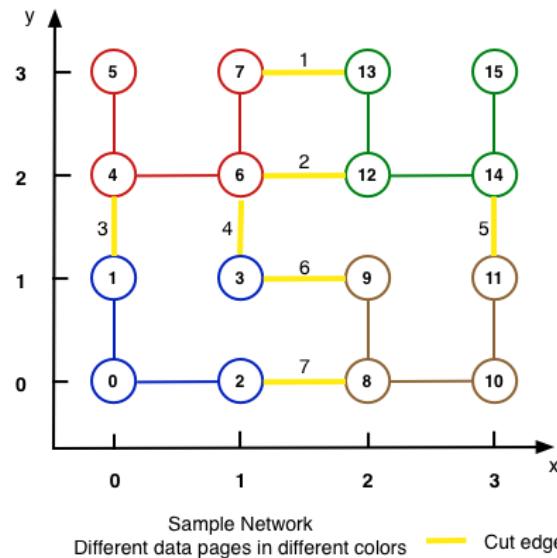
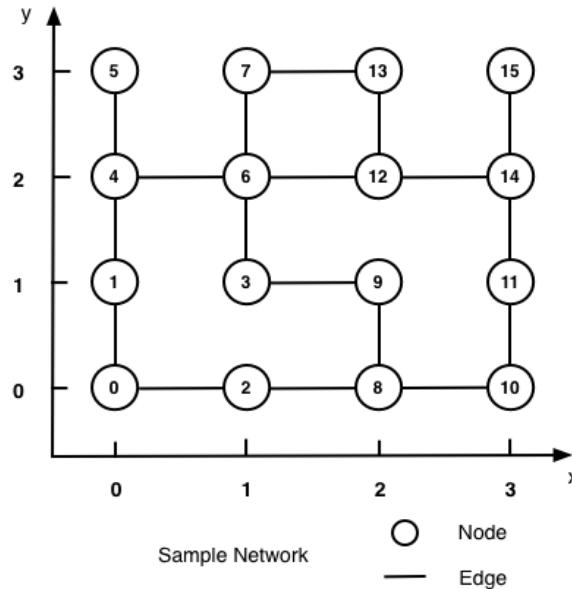
Graph Based Storage Methods

- Insights
 - I/O cost of operations (e.g. get-a-successor) minimized by maximizing CRR
 - $CRR = \Pr.$ (node-pairs connected by an edge are together in a disk sector)

File-Structures:

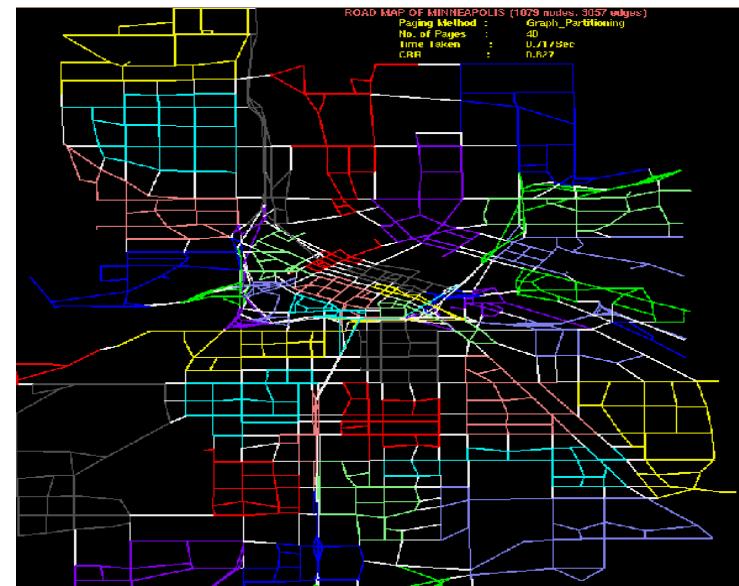
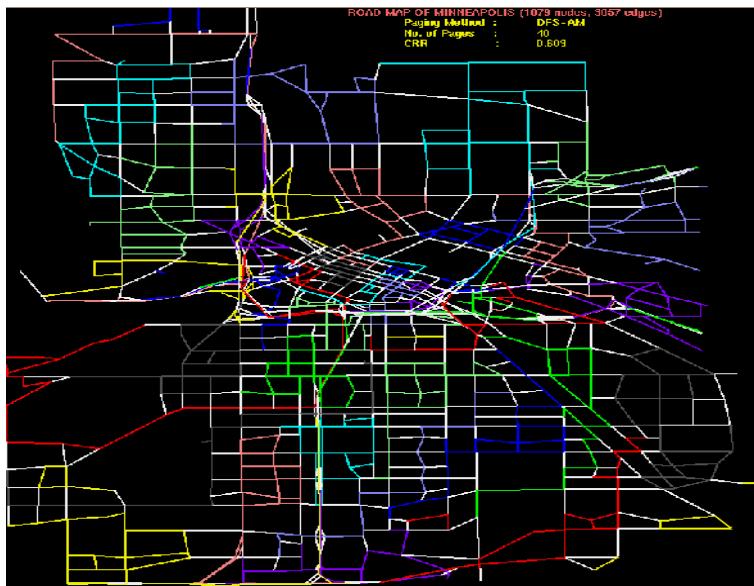
Partition Graph into Disk Blocks

- Which partitioning reduces disk I/O for graph operations?
 - Choice 1: Geometric partition
 - Choice 2: Min-cut Graph Partition
 - Choice 2 cuts fewer edges and is preferred
 - Assuming uniform querying popularity across edges



Graph Based Storage Methods

- Consider two disk-paging of Minneapolis (明尼阿波利斯) major roads
 - Non-white edges => node pair in same page
 - White edge are cut-edges
 - Node partitions on right has fewer cut-edges and is preferred => higher CRR

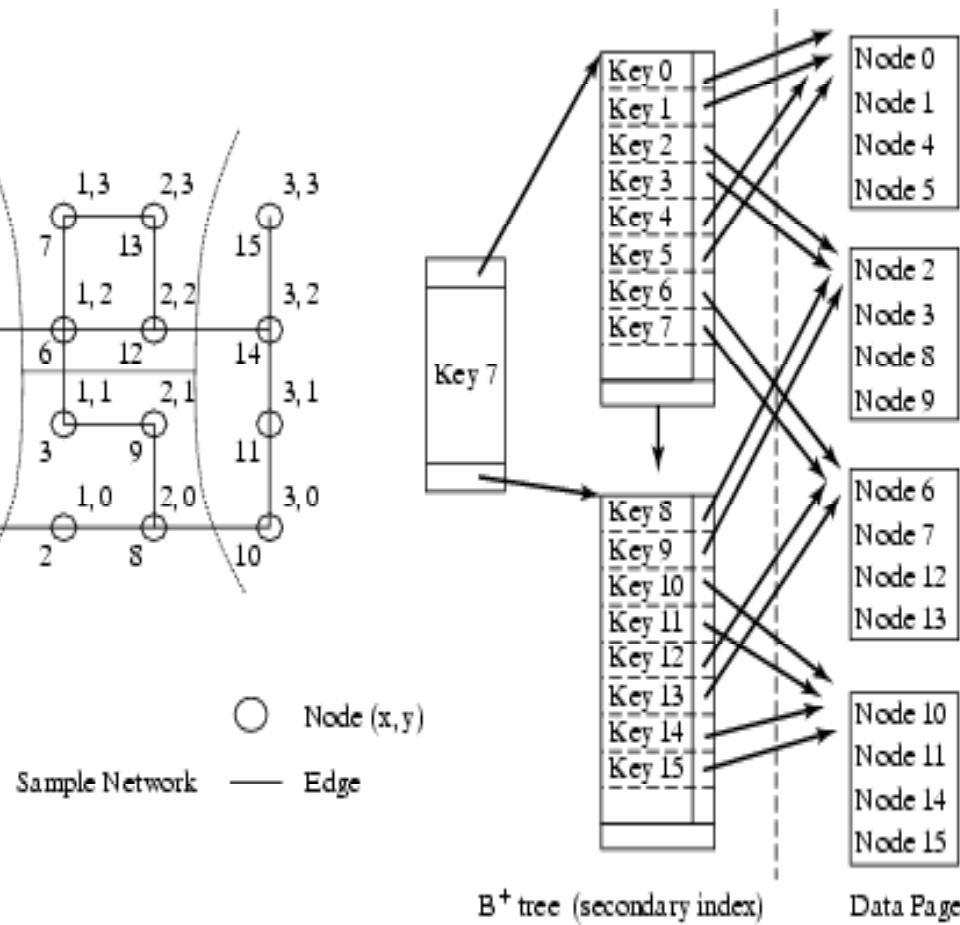
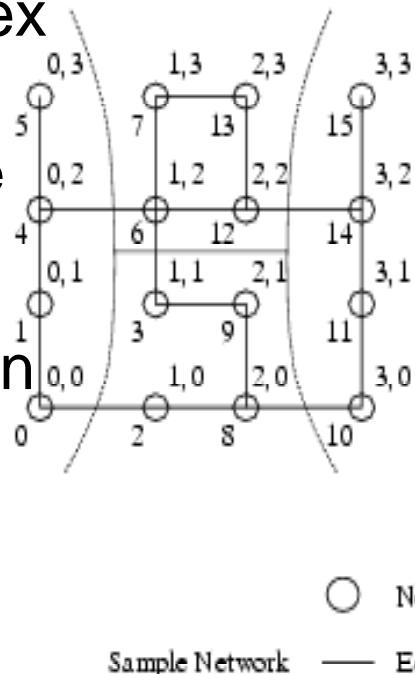


Clustering and Storing a Sample Network

- Storage method idea
 - Divide nodes into sectors
 - to maximize CRR
 - Use a secondary index
 - for find()
 - using R-tree or B-tree

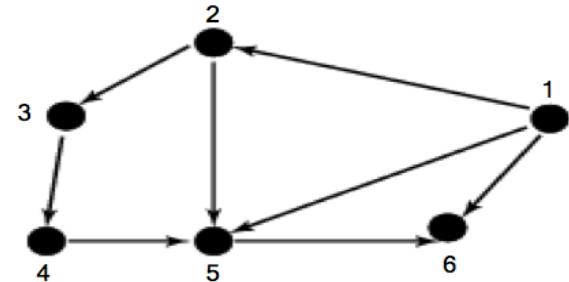
Example

- left part : node division
 - right part
 - disk sectors
 - secondary index
 - B-tree/Z-order



Exercise: Graph Based Storage Methods

- If a disk page holds 3 records, which partitioning will have fewest cut-edges?
 - (1, 2, 3), (4, 5, 6)
 - (2, 3, 4), (1, 5, 6)
 - (1, 2, 6), (3, 4, 5)
 - (1, 3, 5), (2, 4, 6)



Node				
nid	x	y	Successors	Predecessors
1	—	—	(2,5,6)	()
2	—	—	(3,5)	(1)
3	—	—	(4)	(3)
4	—	—	(5)	(3)
5	—	—	(6)	(2,1)
6	—	—	()	(1,5)

Quiz 8

- Which of the following is not disk-based representations of graphs?
 - a) Normalized tables (e.g., node table and edge table)
 - b) Denormalized table (e.g., node table with successors and predecessors columns)
 - c) Adjacency matrix
 - d) All of the above

9.4.2 Query Processing for Spatial Networks

- Query Processing
 - DBMS decomposes a query into building blocks
 - Keeps a couple of strategy for each building block
 - Selects most suitable one for a given situation
- Building blocks
 - Connectivity(A, B): Is node B reachable from node A?
 - Shortest path(A, B): Identify the least cost path from node A to node B

Algorithms

- Main memory (数据结构基础)
 - Connectivity: Breadth first search, Depth first search
 - Shortest path: Dijkstra's algorithm, A*
 - Estimated cost in A* for spatial data: euclidean distance
- Disk-based
 - Shortest path - Hierarchical routing algorithm
 - Connectivity strategies are in SQL3

Quiz 9

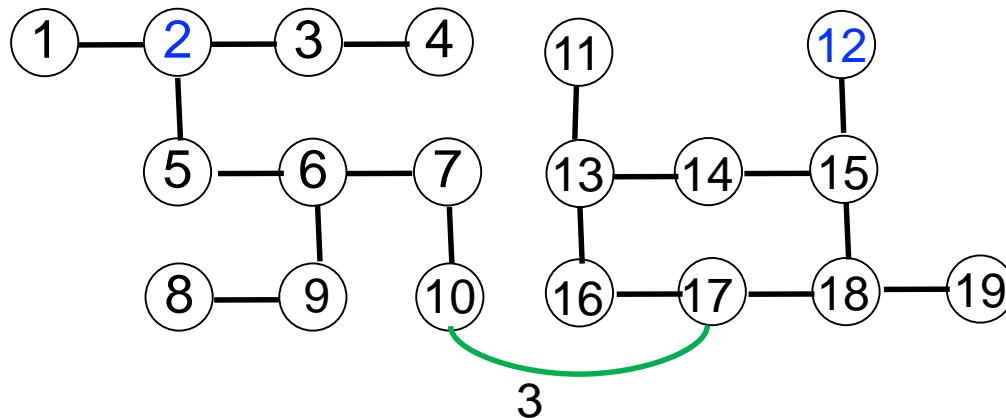
- Which of the following is false?
 - a) Breadth first search visits nodes layer (i.e. generation) by layer
 - b) Depth first search try a path till dead-end, then backtrack to try different paths
 - c) Depth first search always performs better than breadth first search
 - d) None of the above

Shortest Path Strategies

- Dijkstra's and Best first algorithms
 - Work well when entire graph is loaded in main memory
 - Otherwise their performance degrades substantially
- Hierarchical Routing Algorithms
 - Works with graphs on secondary storage
 - Loads small pieces of the graph in main memory
 - Can compute least cost routes

Hierarchical Routing: Simple Example

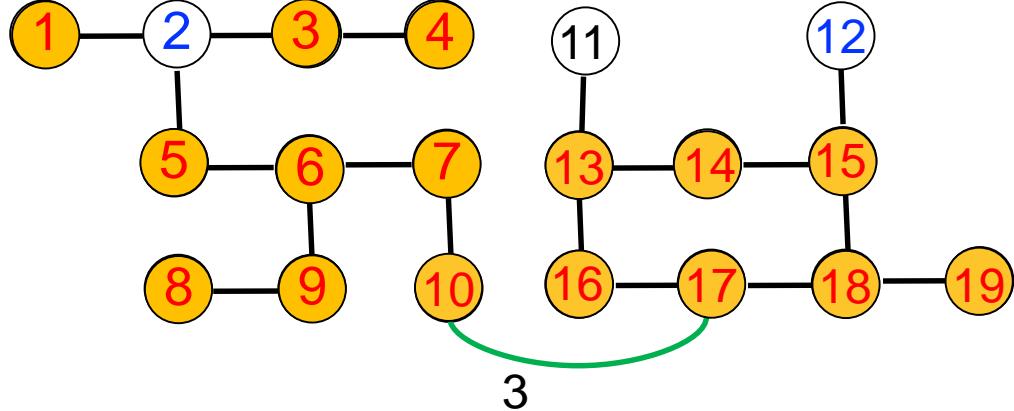
- Goal: Find the shortest path between Nodes 2 and 12
- Candidate algorithms: Dijkstra's algorithm, A*, and hierarchical routing



The edge length is 1 for every edge, except $\text{length}(\text{edge } (11-18)) = 3$

Trace Dijkstra's algorithm

- Tie-breaking: prefer node with higher index



18 nodes expanded

Expanded: Open List:

2	1, 3, 5
5	1, 3, 6
3	1, 4, 6
1	4, 6
6	4, 7, 9
4	7, 9
9	7, 8
7	8, 10
10	8, 17
8	17
17	16, 18
18	15, 16, 19
16	13, 15, 19
19	13, 15
15	12, 13, 14
14	12, 13
13	11, 12
12	

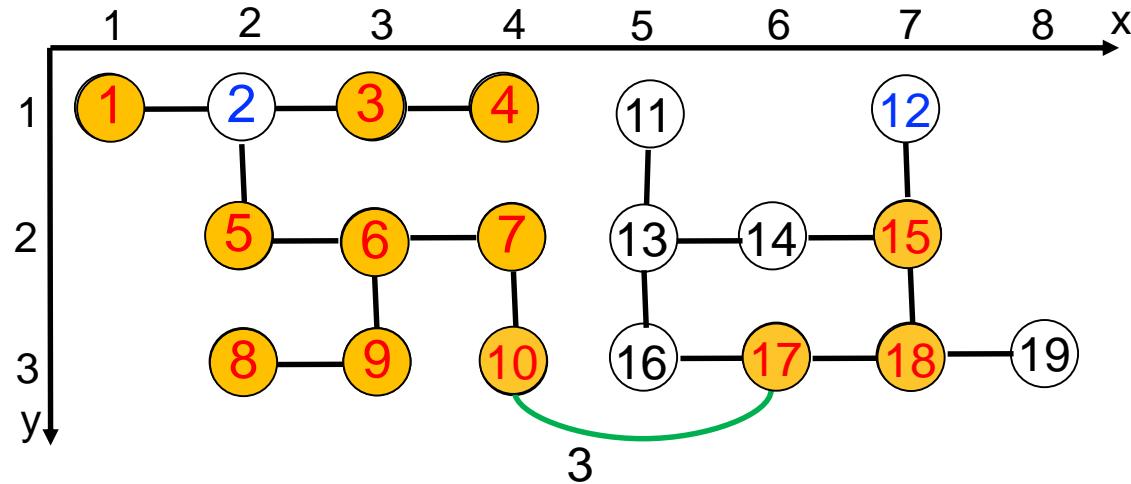
Q? Will A* (with Euclidean distance) be efficient?

Trace for A* algorithm

Cost(node n) = graph_distance (2, n) + Euclidean_distance (n, 12)

14 nodes expanded

Tie-breaking: prefer node with higher index

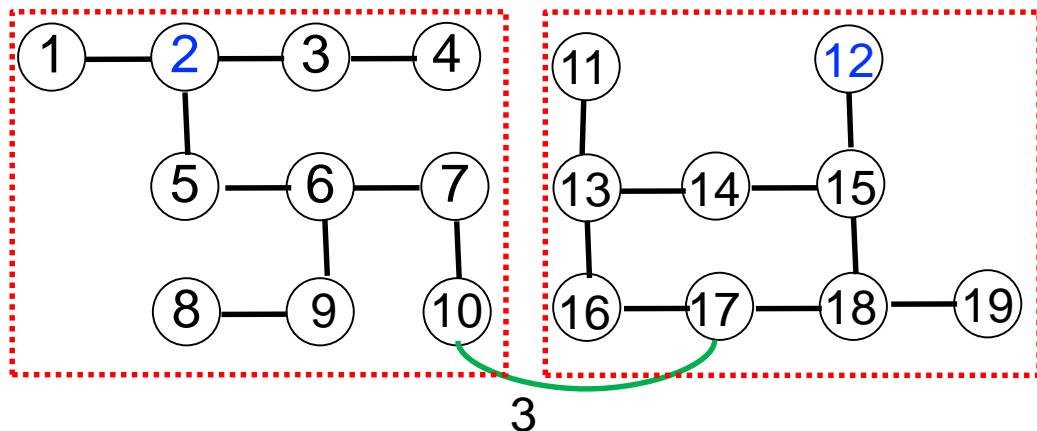


Q? How may we reduce computation cost further?

Expand:	Open List:
2	1, 3, 5
3	1, 4, 5
4	1, 5
5	1, 6
6	1, 7, 9
7	1, 9, 10
10	1, 9, 17
1	9, 17
9	8, 17
17	8, 16, 18
8	16, 18
18	15, 16, 19
15	12, 14, 16, 19
12	

Core Idea of Hierarchical Routing

- Recognize **Islands** and **bridges**
- $\text{SP}(2,12)$ must include **bridge** edge $(10,17)$
- Divide n conquer
$$\text{SP}(2,12) = \text{SP}(2, 10) + \text{edge}(10,17) + \text{SP}(17, 12)$$
- Generalize to the case of multiple **bridges**



Trace Hierarchical Routing

4 nodes expanded in

SP(17 , 12) using A*

Expanded:

17	16,18
18	15,16,19
15	12,14,16,19
12	

Open List:

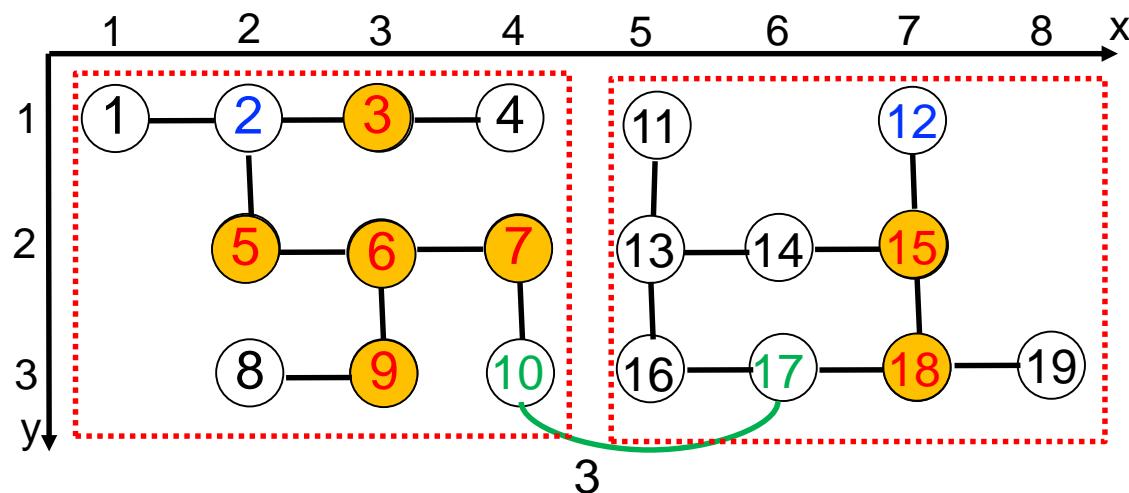
#Nodes expanded: $7 + 4 = 11$

7 nodes expanded in SP(2 , 10) using A*

Expanded:

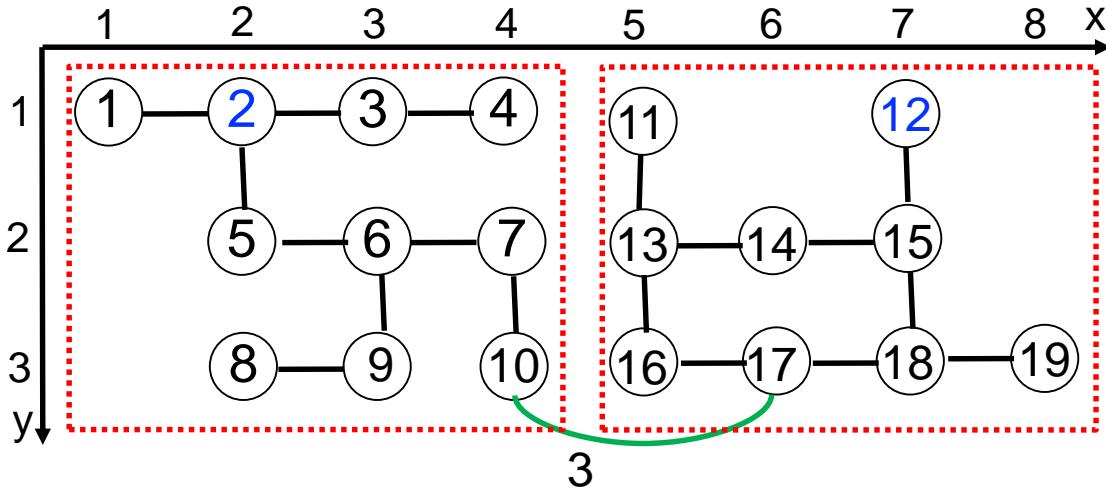
2	1,3,5
5	1,3,6
3	1,4,6
6	1,4,7,9
9	1,4,7,8
7	
10	1,4,8,10

Open List:



Did We Reduce Computational Cost?

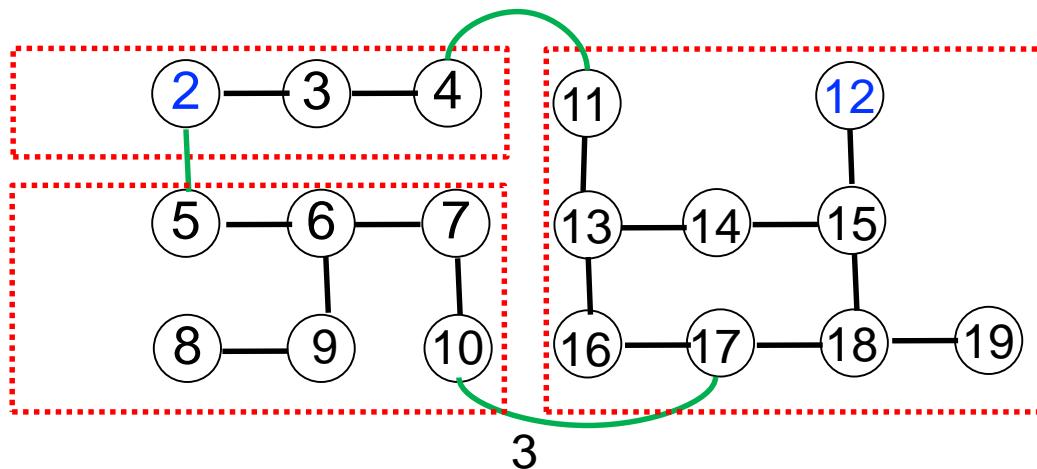
Algorithm (World View)	# of nodes expanded
Dijkstra's (Graphs)	18
A* (Spatial Graphs)	14
Hierarchical Routing (Islands)	$7 + 4 = 11$



Q? What if Multiple bridges?
Q? How to choose among bridges?

Challenges: Multiple Islands & Bridges

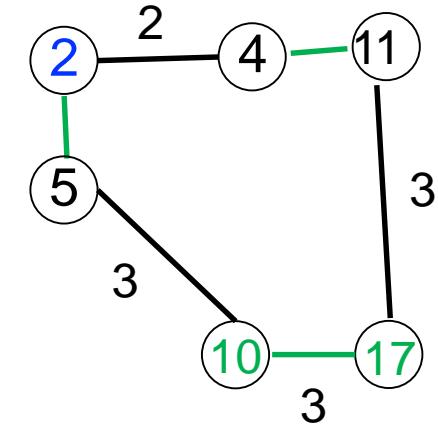
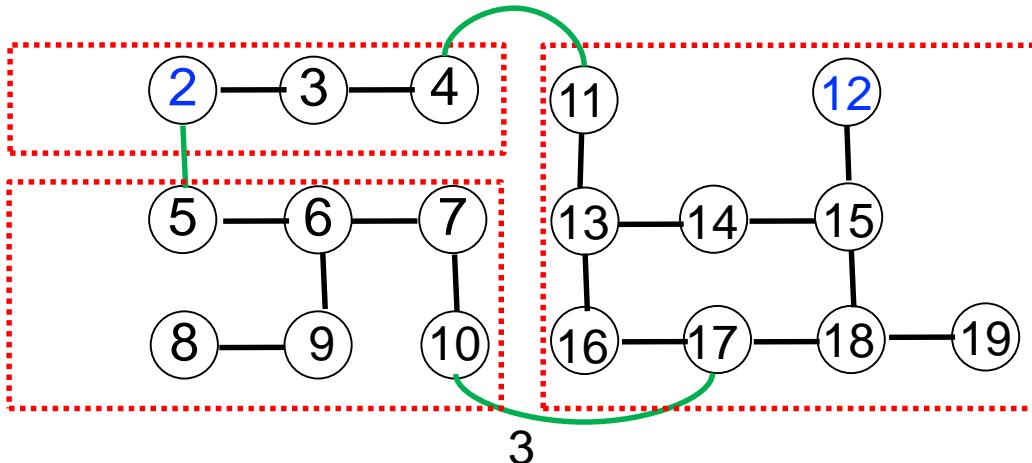
- Invariant: SP(2,12) must use one or more bridges
- Challenge 1: Multiple islands increase computational cost
- Challenge 2: Multiple bridges per island increase computational cost



Hierarchical Algorithm with Multiple Islands

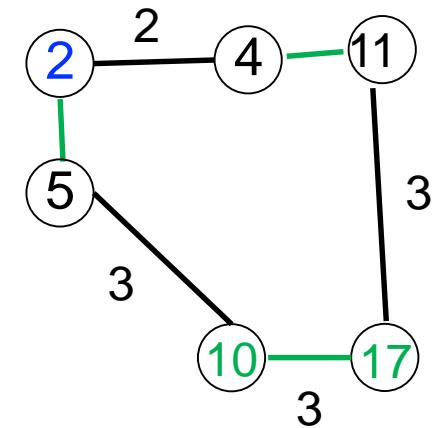
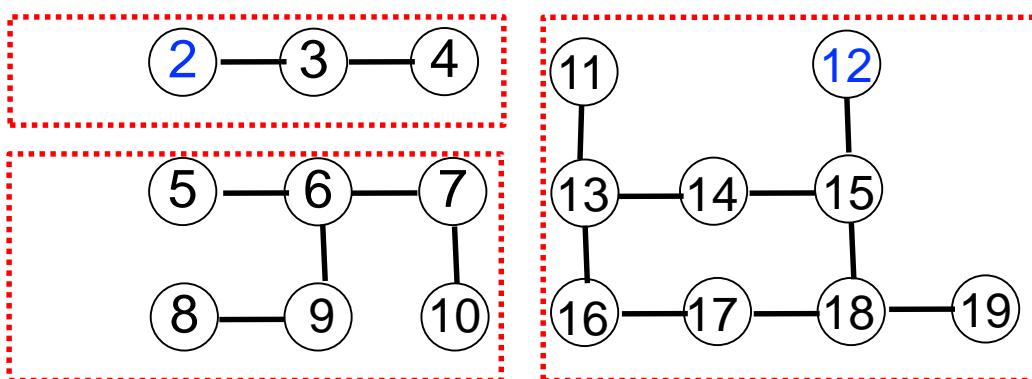
Data Structures:

- Port node (a.k.a., boundary node) : a node with edges to multiple islands
- Port Graph (a.k.a., boundary graph) :
- Island graphs (a.k.a., fragment graphs)
- Precompute & store
 - Shortest_path_costs: SPC(node i, node j) for all (or selected) node pairs (j,j)



Hierarchical Algorithm with Multiple Islands

- Choose port pair (a.k.a., boundary node pair) for SP(2,12)
 - min SPC(2, local port) + SPC(2's local port, 12's local port) + SPC (local port, 12)
 - Choices: (2,11), (2, 17), (4, 11), (4, 17)
 - Chosen port pair: (4, 11)
- Divide and conquer: $SP(2,12) = SP(2, 4) \cdot SP(4, 11) \cdot SP(11, 12)$
 - Use Dijkstra's or A* for sub-problems
- Refine algorithm to reduce storage cost



Hierarchical Routing: Key Ideas

- Key ideas behind Hierarchical Routing Algorithm
 - Fragment graphs - pieces of original graph obtained via node partitioning
 - Boundary nodes - nodes of with edges to two fragments
 - Boundary graph - a summary of original graph
 - Contains boundary nodes
 - Boundary edges: edges across fragments or paths within a fragment

Hierarchical Routing: Insight

- A Summary of optimal path in original graph can be computed
 - Using **boundary graph** and 2 fragments
- The summary can be expanded into optimal path in original graph
 - Examining a fragment overlapping with the path
 - Loading one fragment in memory at a time

Quiz 10

- Which of the following is false?
 - a) Hierarchical routing algorithms are Disk-based shortest path algorithms
 - b) Breadth first search and depth first search are both connectivity query algorithms
 - c) Best first algorithm is always faster than Dijkstra's algorithm
 - d) None of the above

第九章 空间网络模型与查询

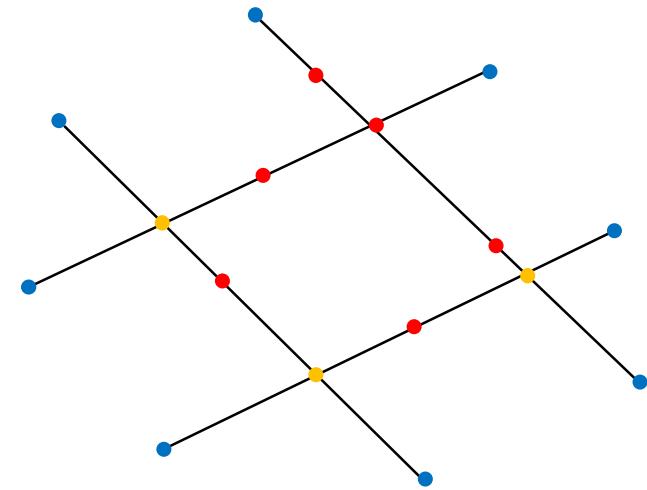
- 9.1 Motivation and use cases
- 9.2 Conceptual model
- 9.3 Logical model
 - 9.3.1 Transitive Closure
 - 9.3.2 CONNECT statement
 - 9.3.3 RECURSIVE statement
 - 9.3.4 RECURSIVE in PostgreSQL
- 9.4 Physical model
 - 9.4.1 Storage and data structures
 - 9.4.2 Algorithms for connectivity query and shortest path
- 9.5 pgRouting

参考教材：

Spatial Databases: A Tour Chapter 6
空间数据库管理系统概论，3.3-3.4

几何对象模型 → 空间网络模型

- 节点选择
 - 所有顶点 X
 - 网络节点过多
 - 线段端点
 - 一条道路仅两个端点
 - 线段交点
 - 增加节点，增强网络连通性
 - 距离相近的顶点合并
 - 减少节点，增强网络连通性
- 边选择
 - 节点之间的道路
 - 有向图、无向图



pgRouting

- pgRouting - geospatial routing functionality
- Advantages of the database routing approach
 - Data and attributes can be modified by many clients, like QGIS through JDBC, ODBC, or directly using PI/pgSQL
 - Data changes can be reflected instantaneously through the routing engine
 - The “cost” parameter can be dynamically calculated through SQL and its value can come from multiple fields or tables
- How to create a database and load pgRouting
 - createdb netdatabase <http://docs.pgrouting.org/>
 - psql netdatabase -c “create extension postgis”
 - psql netdatabase -c “create extension **pgrouting**”

pgRouting

- How to build a topology
 - Ends of an edge will be connected to a unique node
 - Select `pgr_createTopology('myroads', 0.000001)`
 - 线段端点作为网络节点，合并距离相近的节点，距离相近判断阈值为参数`0.000001`，构建网络
- How to check your graph for errors
 - Select `pgr_analyzeGraph('myroads', 0.000001);`
 - Select `pgr_analyzeOneway('myroads', s_in_rules, s_out_rules, t_in_rules, t_out_rules, direction);`
 - 分析当前网络是否存在距离相近的节点，是否存在线段交点

pgRouting

- How to build a topology
 - Select `pgr_createTopology('myroads', 0.000001)`
- How to check your graph for errors
 - Select `pgr_analyzeGraph('myroads', 0.000001);`
- How to renode if there is any error
 - `pgr_nodeNetwork`
 - 将相交的线段按交点分成多条线段
- How to compute a route
 - Select `pgr_<algorithm>` (<SQL for edges>, start, end, <additional options>
 - `pgr_dijkstra(text edges_sql, bigint start_vid, bigint end_vid, boolean directed:=true)`
 - `pgr_costResult[]` and `pgr_geomResult[]`

pgRouting Data Types

- pgr_costResult[]
 - `seq` integer -- sequential ID indicating the path order
 - `id1` integer -- typically the node id
 - `id2` integer -- typically the edge id
 - `cost` float8 -- cost attribute
- pgr_geomResult[]
 - `seq` integer -- sequential ID indicating the path order
 - `id1` integer -- generic name, to be specified by the func
 - `id2` integer -- generic name, to be specified by the func
 - `geom` geometry

pgRouting Topology Functions

- Topology of a network
 - An edge table with source and target attributes
 - A vertices table associated with the edge table
- pgr_createTopology
 - Build a network topology based on the geometry information
 - Parameters
 - `edge_table` text - Newtork table name
 - `tolerance` float8 - Snapping tolerance of disconnected edge
 - `the_geom` text - Geometry colum name of the network table
 - `id` text - Primary key column name of the network table
 - `source` text - Source column name of the network table
 - `target` text - Target column name of the network table
 - `rows_where` text - Condition to SELECT a subset or rows

pgRouting Topology Functions

- pgr_createTopology
 - Builds a network topology based on the geometry information
 - The edge_table will be affected
 - The `source` column values will change
 - The `target` column values will change
 - An `index` will be created, if it doesn't exists, to speed up the process to the following columns
 - id, the_geom, source, target
 - OK after the network topology has been built
 - Creates a vertices table: `<edge_table>_vertices_pgr`
 - Fills `id` and `the_geom` columns of the vertices table
 - Fills the `source` and `target` columns of the edge table referencing the id of the vertices table

pgRouting Topology Functions

- pgr_createTopology
 - Builds a network topology based on the geometry information
 - The edge_table will be affected
 - OK after the network topology has been built
 - FAIL when the network topology was not built due to an error
 - A required column of the Network table is not found or is not of the appropriate type
 - The condition is not well formed
 - The names of source, target or id are the same
 - The SRID of the geometry could not be determined

pgRouting Topology Functions

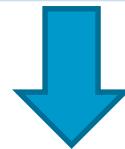
- pgr_createTopology
 - The Vertices Table
 - `id` bigint - Identifier of the vertex.
 - `cnt` integer - Number of vertices in the edge_table that reference this vertex
 - `chk` integer - Indicator that the vertex might have a problem
 - `ein` integer - Number of vertices in the edge_table that reference this vertex AS incoming
 - `eout` integer - Number of vertices in the edge_table that reference this vertex AS outgoing
 - `the_geom` geometry - Point geometry of the vertex
- prg_createVerticesTable
 - Reconstructs the vertices table based on the source and target information

pgRouting Topology Functions

id	the_geom	source	target	...
81	(0 20, 10 20.01)			
45	(10 19.99, 10 2)			
...	...			

几何表
edge_table

边表 edge_table



pgr_createTopology
tolerance =0.05

id	the_geom	source	target	...
81	(0 20, 10 20.01)	3	5	
45	(10 19.99, 10 2)	5	7	
...	

顶点表 edge_table_vertices_pgr

id	cnt	chk	ein	eout	the_geom
3					(0 20)
5					(10 20)
...					...

pgRouting Topology Functions

- pgr_analyzeGraph
 - Analyze the network topology
 - The edge table to be analyzed must contain
 - a **source** column and a **target** column filled with id's of the vertices of the segments
 - the corresponding vertices table **<edge_table>_vertices_pgr** that stores the vertices information

```
varchar pgr_analyzeGraph(text edge_table, double precision tolerance,
                          text the_geom:='the_geom', text id:='id',
                          text source:='source', text target:='target', text rows_where:='true')
```

- pgr_analyzeOneway
 - Analyzes oneway streets and identifies flipped segments

pgRouting Topology Functions

- pgr_analyzeGraph
 - Analyzes the network topology
 - `edge_table`: text Network table name. (may contain the schema name as well)
 - `tolerance`: float8 Snapping tolerance of disconnected edges. (in projection unit)
 - `the_geom`: text Geometry column name of the network table. Default value is the_geom
 - `id`: text Primary key column name of the network table. Default value is id
 - `source`: text Source column name of the network table. Default value is source
 - `target`: text Target column name of the network table. Default value is target
 - `rows_where`: text Condition to select a subset or rows. Default value is true to indicate all rows

pgRouting Topology Functions

- pgr_analyzeGraph
 - Analyzes the network topology
 - OK after the analysis has finished
 - Uses the vertices table: `<edge_table>_vertices_pgr`
 - Fills completely the `cnt` and `chk` columns of the vertices table
 - Returns the analysis of the section of the network defined by `rows_where`
 - FAIL when the analysis was not completed due to an error
 - The vertices table is not found
 - A required column of the Network table is not found or is not of the appropriate type
 - The condition is not well formed
 - The names of source, target or id are the same
 - The SRID of the geometry could not be determined

pgRouting Topology Functions

边表 edge_table

id	the_geom	source	target	...
81	(0 20, 10 20)	3	5	
45	(10 20, 10 2)	5	7	
...	

顶点表 edge_table_vertices_pgr

id	cnt	chk	ein	eout	the_geom
3					(0 20)
5					(10 20)
...					..

pgr_analyzeGraph

tolerance =0.05

顶点表 edge_table_vertices_pgr

id	cnt	chk	ein	eout	the_geom
3	2	0			(0 20)
5	4	1			(10 20)
...

pgRouting Topology Functions

- pgr_nodeNetwork
 - Node an network edge table
 - This function reads the `edge_table` table, that has a primary key column `id` and geometry column named `the_geom` and `intersect` all the segments in it against all the other segments and then `creates` a table `edge_table_noded`
 - It uses the `tolerance` for deciding that multiple nodes within the tolerance are considered `the same node`

pgRouting Topology Functions

- pgr_nodeNetwork
 - Node an network edge table
 - `edge_table`: text Network table name. (may contain the schema name as well)
 - `tolerance`: float8 tolerance for coincident points (in projection unit)
 - `id`: text Primary key column name of the network table. Default value is id
 - `the_geom`: text Geometry column name of the network table. Default value is the_geom
 - `table_ending`: text Suffix for the new table's. Default value is noded

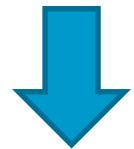
pgRouting Topology Functions

- pgr_nodeNetwork
 - Nodes an network edge table
 - The output table will have for edge_table_noded
 - `id`: bigint Unique identifier for the table
 - `old_id`: bigint Identifier of the edge in original table
 - `sub_id`: integer Segment number of the original edge
 - `source`: integer **Empty** source column to be used with `pgr_createTopology` function
 - `target`: integer **Empty** target column to be used with `pgr_createTopology` function
 - `the_geom`: geometry Geometry column of the noded network

pgRouting Topology Functions

边表 edge_table

id	the_geom	source	target	...
81	(0 20, 10 20)			
45	(10 20, 10 2)			
...	...			

 pgr_nodeNetwork
tolerance =0.05

边表 edge_table_noded

id	oid	subid	source	target	the_geom
34	81	1			(0 20, 5 20)
35	81	2			(5 20, 10 20)
36	45	1			(10 20, 10 2)
...
...

pgRouting Routing Functions

- All Pairs Shortest Path, Johnson's Algorithm
- All Pairs Shortest Path, Floyd-Warshall Algorithm
- Shortest Path A*
- Bi-directional Dijkstra Shortest Path
- Bi-directional A* Shortest Path
- Shortest Path Dijkstra
- Driving Distance
- K-Shortest Path, Multiple Alternative Paths
- K-Dijkstra, One to Many Shortest Path
- Traveling Sales Person <http://docs.pgrouting.org/>
- Turn Restriction Shortest Path (TRSP)

Dijkstra最短路径算法

- `pgr_dijkstra(text edges_sql, bigint start_vid, bigint end_vid, boolean directed:=true)`
 - `edges_sql`: a SQL query
 - select id, source, target, cost [, reverse_cost] from edge_table
 - `id`: ANY-INTEGER identifier of the edge
 - `source`: ANY-INTEGER identifier of the first end point vertex of the edge
 - `target`: ANY-INTEGER identifier of the second end point vertex of the edge
 - `cost`: ANY-NUMERICAL weight of the edge (source, target), if negative: edge (source, target) does not exist, therefore it's not part of the graph
 - `reverse_cost`: ANY-NUMERICAL (optional) weight of the edge (target, source), if negative: edge (target, source) does not exist, therefore it's not part of the graph

Dijkstra最短路径算法

- pgr_dijkstra(**text** edges_sql, **bigint** start_vid, **bigint** end_vid, **boolean** directed:=true)
 - **edges_sql**: a SQL query
 - select id, source, target, cost [, reverse_cost] from edge_table
 - **start_vid**: BIGINT identifier of the starting vertex of the path
 - **end_vid**: BIGINT identifier of the ending vertex of the path
 - **directed**: boolean (optional). When false the graph is considered as Undirected. Default is true which considers the graph as Directed

Dijkstra最短路径算法

- pgr_costResult[] pgr_dijkstra(text sql, integer source, integer target, boolean directed:=true)

```
SELECT * FROM pgr_dijkstra(
    'SELECT id, source, target, cost, reverse_cost FROM edge_table',
    2, 3
);
seq | path_seq | node | edge | cost | agg_cost
----+-----+-----+-----+-----+
  1 |      1 |    2 |    4 |    1 |      0
  2 |      2 |    5 |    8 |    1 |      1
  3 |      3 |    6 |    9 |    1 |      2
  4 |      4 |    9 |   16 |    1 |      3
  5 |      5 |    4 |    3 |    1 |      4
  6 |      6 |    3 |   -1 |    0 |      5
(6 rows)

SELECT * FROM pgr_dijkstra(
    'SELECT id, source, target, cost, reverse_cost FROM edge_table',
    2, 5
);
seq | path_seq | node | edge | cost | agg_cost
----+-----+-----+-----+-----+
  1 |      1 |    2 |    4 |    1 |      0
  2 |      2 |    5 |   -1 |    0 |      1
(2 rows)
```

Dijkstra最短路径算法

- `SELECT seq, id1 AS node, id2 AS edge, cost
FROM pgr_dijkstra(`

```
'SELECT id, source, target, len as cost FROM road_network'
```

5, 3, false

);

- 通过SQL语句构建查询表， 提供网络关系(`id, source, target, cost`)
- 查询约束， 通过`where`语句查询不同网络关系实现
 - 某条道路在修路: `where id != 101`
 - 只走`cost`小于10的道路: `where cost < 10`
 -

Example

- create table Road (
 id serial primary key,
 name text,
 geom geometry(LineString, 4326));
- insert into Road(name, geom) values
 ('A', ST_GeomFromText('LineString(0 20, 10 20)', 4326));
- ('B', ST_GeomFromText('LineString(10 20, 10 2)', 4326));
- ('C', ST_GeomFromText('LineString(10 2, 20 2)', 4326));
- ('D', ST_GeomFromText('LineString(0 5, 20 5)', 4326));
- ('E', ST_GeomFromText('LineString(20 4.9999, 20 2.0001)', 4326));

Example

- `create table Road_network (`
 `id serial primary key,`
 `name text,`
 `source int,`
 `target int,`
 `geom geometry(LineString, 4326),`
 `len float);`
- `insert into Road_network(name, geom) select name, geom from Road;`
- `select pgr_createTopology('road_network', 0.00001, 'geom', 'id',`
`'source', 'target', 'true');`
- `select pgr_analyzeGraph('road_network', 0.00001, 'geom', 'id', 'source',`
`'target', 'true');`

Example

- `select pgr_nodeNetwork('road_network', 0.00001, the_geom:='geom', id:='id', table_ending:='1');`
- `select pgr_createTopology('road_network_1', 0.001, 'geom', 'id', 'source', 'target', 'true');`
- `select pgr_analyzeGraph('road_network_1', 0.001, 'geom', 'id', 'source', 'target', 'true');`

Example

- create table temp (
 id serial primary key,
 name text,
 source int,
 target int,
 geom geometry(LineString, 4326));
- insert into temp(name, source, target, geom)
 select O.name, N.source, N.target, N.geom from Road_network O,
Road_network_1 N where O.id = N.old_id order by O.name;
- delete from Road_network;
- insert into Road_network(name, source, target, geom)
 select name, source, target, geom from temp;

Example

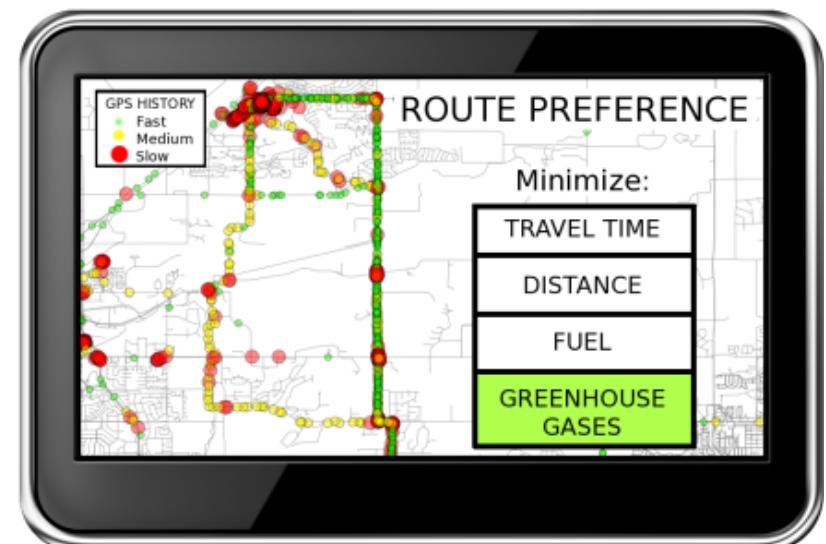
- update Road_network
set len = ST_LENGTH(geom);
- SELECT seq, id1 AS node, id2 AS edge, cost
FROM pgr_dijkstra(
 'SELECT id, source, target, len as cost FROM Road_network',
 5, 3, false
);

Example

- 扩展
 - 新增加路F, 从(0,5)到(10,20)
 - 查询(12,6)到(10,20)的最短距离
 - 查询(12,6)到(18,4)的最短距离
 - 限制单行, 只能从(10,5)到(0,5)

Trends: Next Generation Navigation Services

- Eco-Routing
 - Energy cost may be negative, e.g., downhill
 - Dijkstra's assumptions
 - Sub-path optimality
 - Fixed, non-negative and additive edge costs
 - Relaxing assumptions
 - Fixed and additive edge cost (maybe negative)
 - No “negative cycle”
- Spatio-temporal Graphs
 - Time variant networks
 - Flow networks



Summary

- Spatial Networks are a fast growing applications of SDBs
- Spatial Networks are modeled as graphs
- Graph queries, like shortest path, are transitive closure
 - Not supported in relational algebra
 - SQL features for transitive closure: CONNECT BY, WITH RECURSIVE, pgRouting
- Graph Query Processing
 - Building blocks - connectivity, shortest paths
 - Strategies - Best first, Dijkstra's and Hierarchical routing
- Storage and access methods

空间数据模型总结

- 几何对象模型
- 几何拓扑模型
- 空间网络模型
- 栅格数据模型 (遥感数字图像处理)
- 注记文字模型

